

Introduction to Data Science (Khoa học dữ liệu)

Nguyen Thi Oanh

Hanoi University of Science and Technology
oanhnt@soict.hust.edu.vn

IT4142E, SOICT, HUST, 2018

Contents of the course

- Introduction to Data Science
- Data storage and processing
- Exploratory data analysis
- Machine Learning
- Big data analysis
- Visualization
- Natural language processing
- **Computer vision**
- Graph analysis
- Recommender system

Plan

- Computer Vision and Applications
- Digital image / video ?
- Basic informations about digital images
 - histogram, brightness, contrast, color, texture, ...
 - Library : Opencv
- Convolution and Filters
 - noisy remove,
 - edge detectors
- Feature extraction: local and global descriptor
 - Segmentation
- [Shape recognition]
- [NN is not CV and vice versa]

Computer Vision ?

- Image Processing

- Work with image as a matrix
- Input: image → output: image
- Help human to examine / modify images

- Computer Vision

- Make computers understand images and video
- Images and video are a source of information on the reality



What kind of scene?

Where are the cars?

How far is the building?

...

Computer Vision and Applications

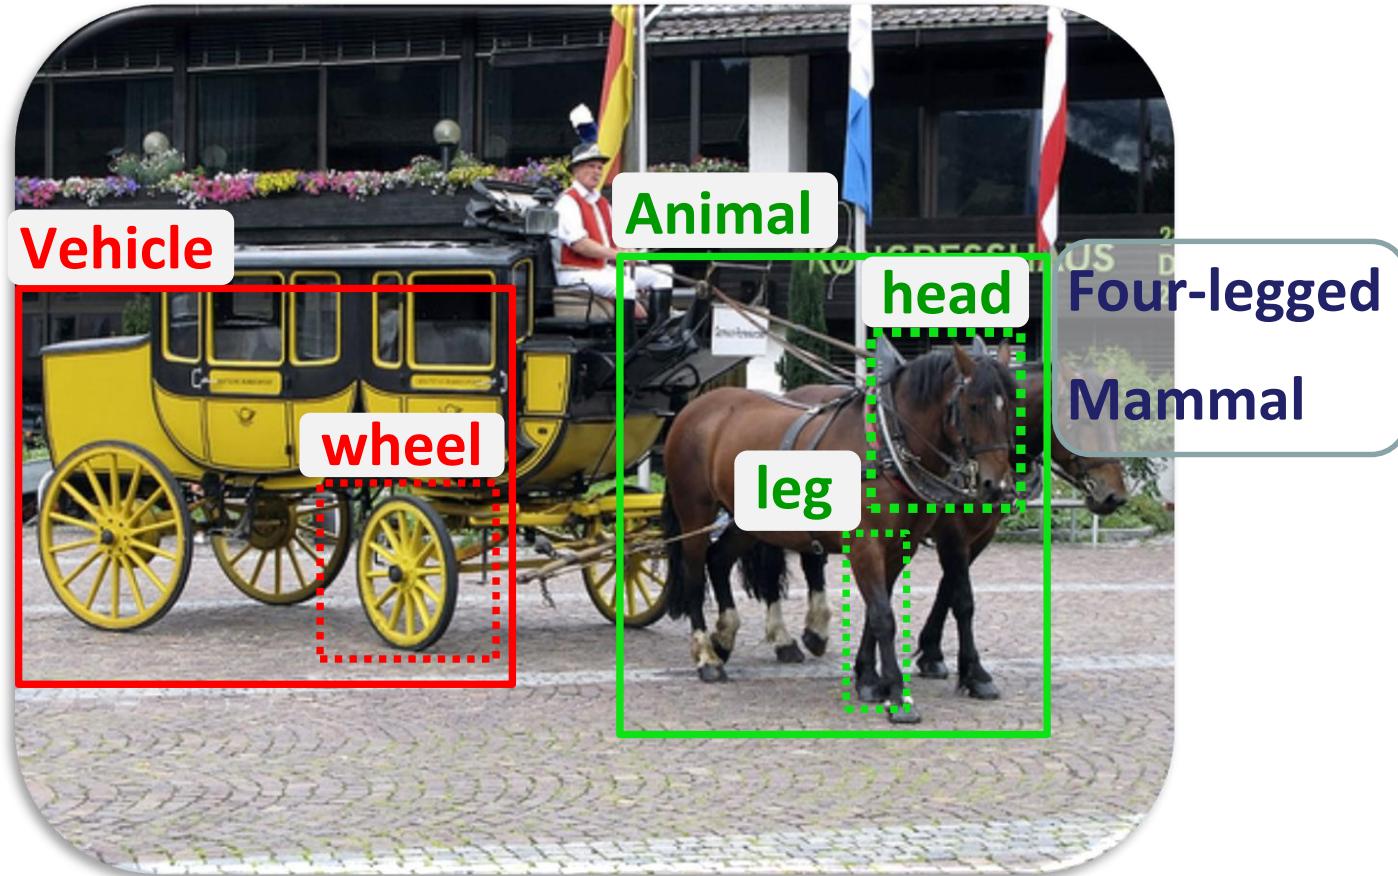
- Images, video are everywhere
- Video, images:
 - Rich information

→ Hot topic, especially
When we talk every day
about AI with smart city,
smart home, smart ...



How vision is used now

- Understand the image



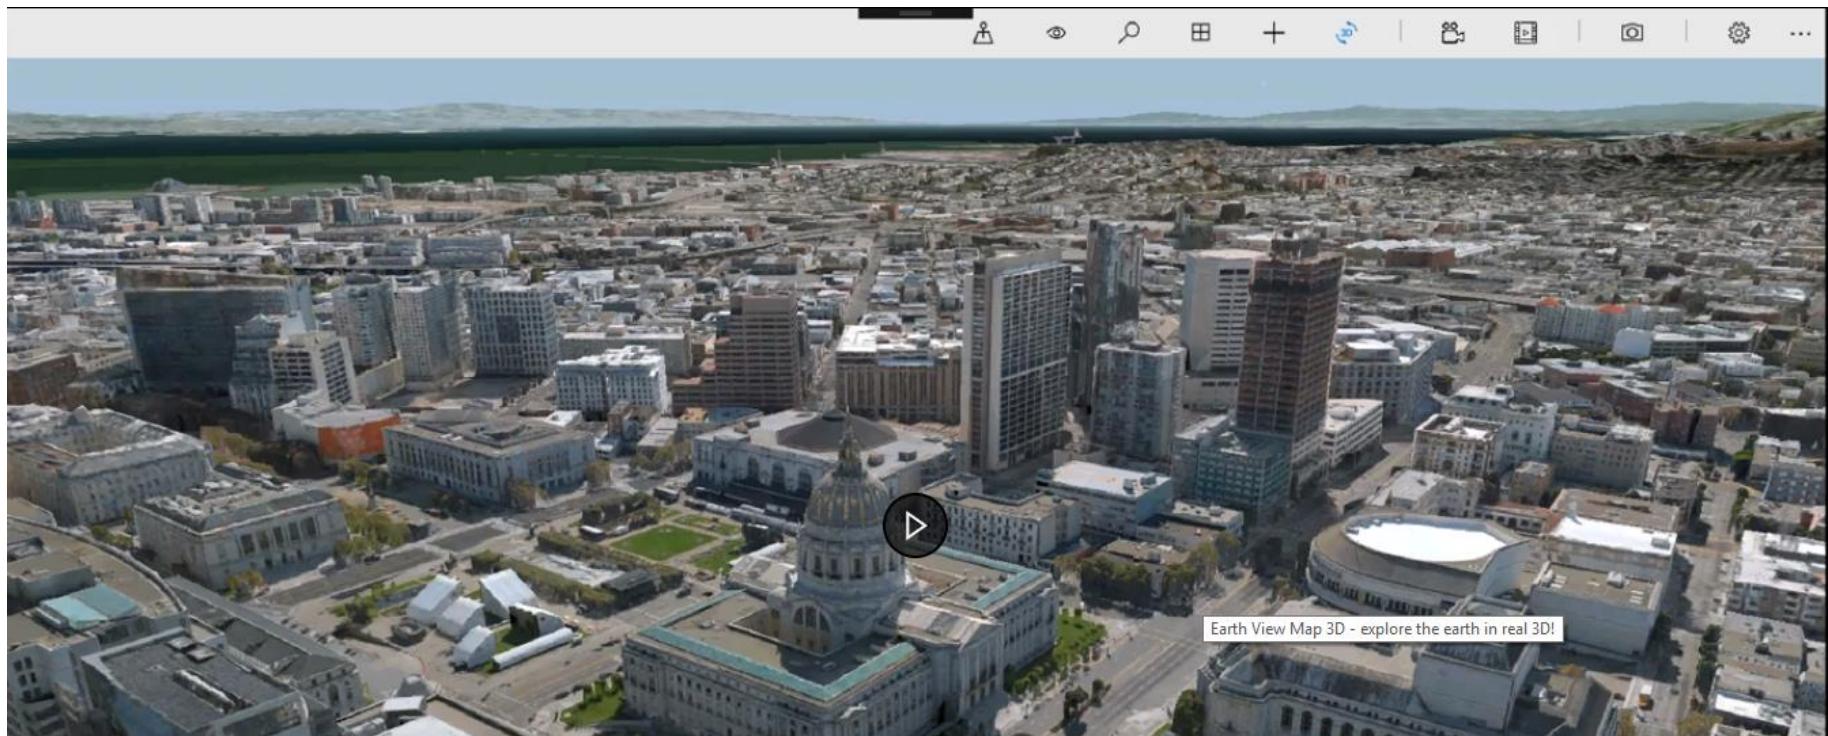
How vision is used now



Facebook's suggestion

How vision is used now

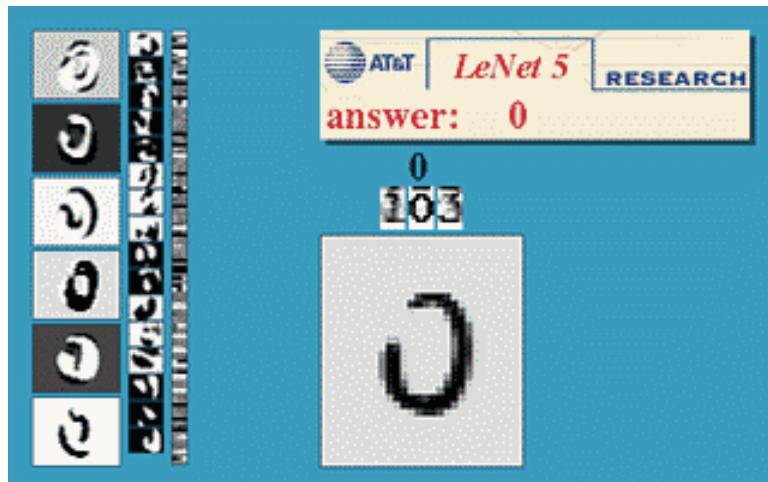
- Earth View, Google earth (3D modeling from lots of 2D images): automatic building generation + hand modeled buildings (Golden Gate bridge or Sydney Opera house)



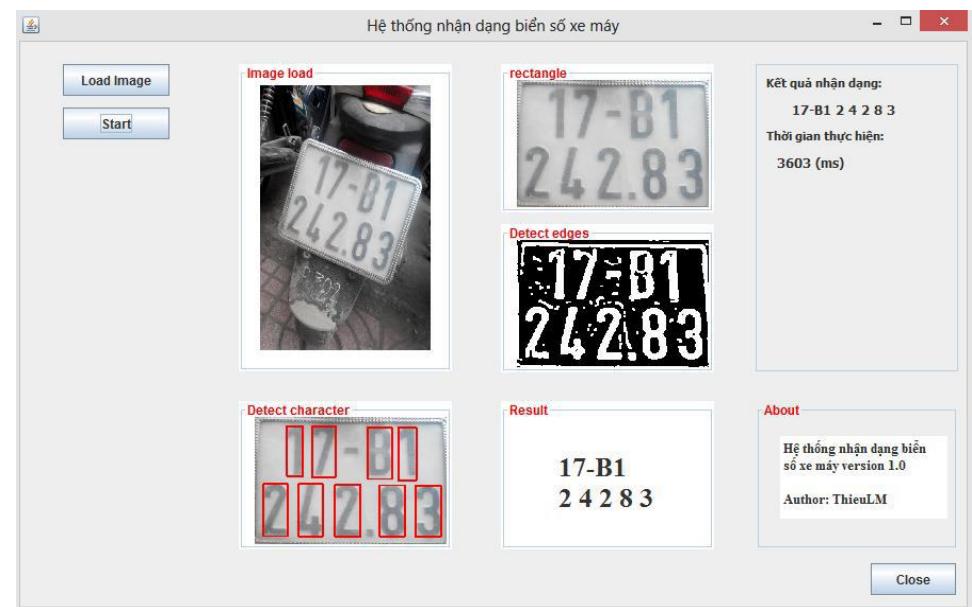
Microsoft's Virtual Earth

How vision is used now

- OCR (Optical character recognition)
- Technology to convert scanned docs to text: each scanner came with an OCR software



Digit recognition, AT&T labs
<http://yann.lecun.com/exdb/lenet/>



Licence plate detection
and character recognition

How vision is used now

- Face detection: most of camera detect faces
 - Main raison is focus, enables smart cropping



How vision is used now

- Smile detection: smart camera
 - Camera can automatically trip the shutter at the right instant to catch the perfect expression

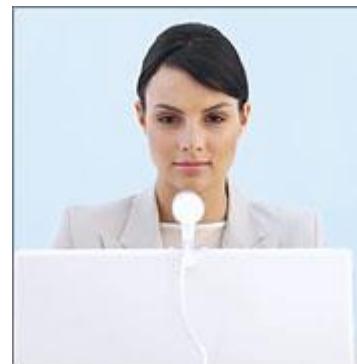


How vision is used now

- Login without a password, but with biometrics (fingerprint, iris, face,..)



Fingerprint scanners on many new laptops,
other devices



Face recognition systems now
beginning to appear more widely
<http://www.sensiblevision.com/>

How vision is used now

- Object recognition (on mobile phones)



Point & Find, Nokia
Google Goggles

How vision is used now

Content-based image retrieval

Google  steve jobs iphone

Search

Web
Images
Maps
Videos
News
Shopping
More

Search by image
Visually similar
More sizes

Any time
Past hour
Past 24 hours
Past week
Past month
Past year
Custom range...



Image size:
620 x 388

Find other sizes of this image:
All sizes - Small - Medium

Best guess for this image: [steve jobs iphone](#)

[The iPhone 5 Suggests That Without Steve Jobs, Apple Is Becomin...](#)

www.forbes.com/sites/.../the-iphone-5-and-the-post-steve-jobs-apple/

3 days ago – see photosClick for full photo gallery: Apple iPhone 5 Event Tim Cook has just wrapped up his introduction of the iPhone 5. On paper, the ...

["Boring" iPhone 5 Is Not A Steve Jobs "Legacy Device" But Cements ...](#)

www.forbes.com/.../boring-iphone-5-is-not-a-steve-jobs-legacy-devi...

2 days ago – see photosGetty ImagesClick for full photo gallery: Apple Introduces The iPhone 5 Wired called the iPhone 5 "utterly boring". The BBC ran a ...

[Visually similar images](#) - Report images



22 Results

Searched over [2.1809 billion](#) images in 1.

for file: test_st_2133673b.jpg

- These results expire in **72 hours**. [Why?](#)
- [Share a success story!](#)
- TinEye is [free](#) to use for non-commercial purposes.

Image Collection Results ([info...](#))



[pa.photoshelter.com](#)
[pa.photoshelter.com/image/I...](#)



[Compare](#) | [Link](#)
JPEG Image
6x672, 393.8 KB

[View 1 more collection result](#)



[blogs.telegraph.co.uk](#)
[steve jobs.jpg](#)
[blogs.telegraph.co.uk/technology/tech...](#)
[blogs.telegraph.co.uk/technology/tech...](#)



[Compare](#) | [Link](#)
JPEG Image
6x200, 78.1 KB

[View all 15 matches](#)



[tecnologia.ig.com.br](#)
[7571361.steve jobs 225_300.jpg](#)
[tecnologia.ig.com.br/noticia/2010/04/...](#)

How vision is used now

- Smart cars → autonomous vehicles

The screenshot shows the Mobileye website's homepage. At the top, there are tabs for "manufacturer products" and "consumer products". A central banner features the slogan "Our Vision. Your Safety." above an overhead view of a car with three cameras labeled: "rear looking camera", "forward looking camera", and "side looking camera". Below this, there are three main sections: "EyeQ Vision on a Chip" (with an image of a chip), "Vision Applications" (with an image of a pedestrian crossing a street), and "AWS Advance Warning System" (with an image of a display screen). To the right, there are two vertical columns: "News" (listing articles like "Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System" and "Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end") and "Events" (listing events like "Mobileye at Equip Auto, Paris, France" and "Mobileye at SEMA, Las Vegas, NV").

Mobileye: vision systems currently in many cars

"In mid 2010 Mobileye will launch a world's first application of full emergency braking for collision mitigation for pedestrians where vision is the key technology for detecting pedestrians"

How vision is used now

- Panorama stitching:



How vision is used now

- Games / robots:



Vision-based interaction game
(Microsoft's Kinect)



<http://www.robocup.org/>



Robot vacuum cleaner

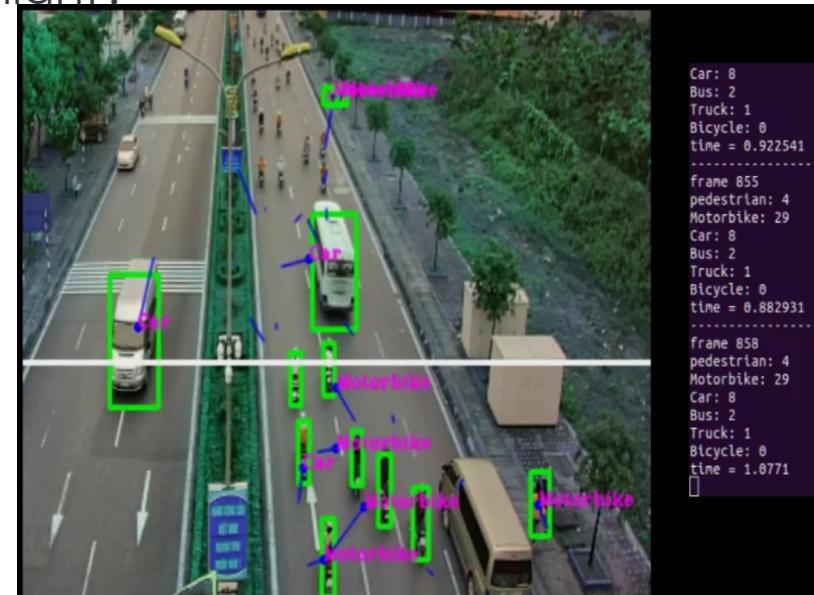
How vision is used now

- To learn more about vision applications and companies, please visit D.Lowe's note:

<https://www.cs.ubc.ca/~lowe/vision.html>

Topics

- Surveillance Cameras:
 - Counting number of clients in shop/restaurant
 - Detecting abnormal behaviors
 - Measuring customer's satisfaction
 - Object tracking: Someone ran a red light?
- Object classification / detection
 - Face / body/ eyes detection
 - Action recognition
 - Defect detection
- Image annotation / labeling
- Text detection and recognition
 - Card visit reader
- 3D object construction from 2D images



What we will talk about?

- 2 Types information we would like to extract from images:
 - Matrix 3D information
 - Semantic Information



How to represent
the content of images ?

3D building?

Where is it?
Text in the picture,
what does it means?
Are there person
in the picture?

Pre-processing

Feature extraction

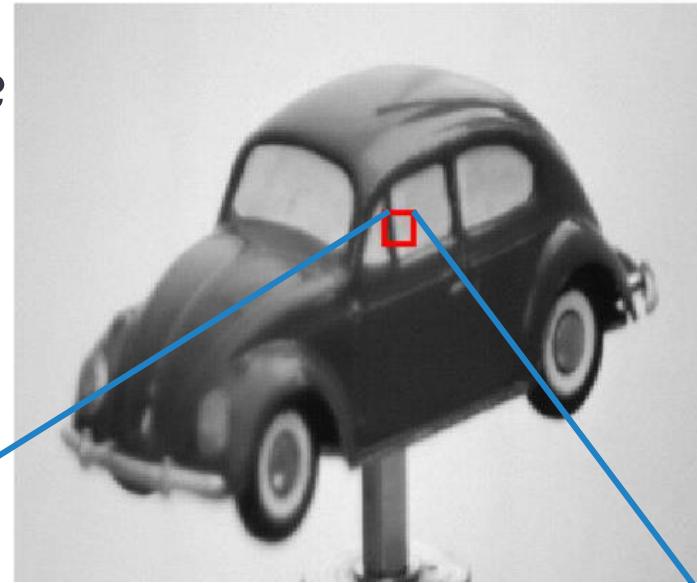
Learning

Deep Learning

Model

Digital images ?

- What we can see on the picture?
 - A car?
- What does the machine see?
 - Image is a matrix of pixels
 - Image N x M : N xM matrix
 - 1 pixel (gray levels):
 - A intensity value: 0-255
 - Black: 0
 - White: 255

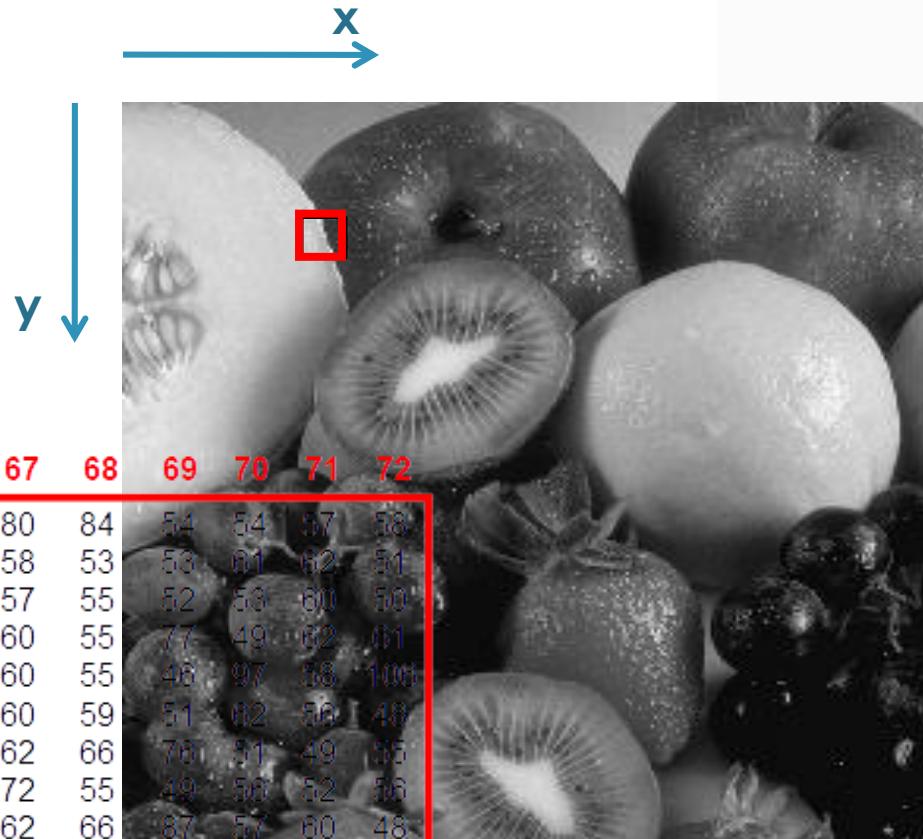


64	60	69	100	149	151	176	182	179
65	62	68	97	145	148	175	183	181
65	66	70	95	142	146	176	185	184
66	66	68	90	135	140	172	184	184
66	64	64	84	129	134	168	181	182
59	63	62	88	130	128	166	185	180
60	62	60	85	127	125	163	183	178
62	62	58	81	122	120	160	181	176
63	64	58	78	118	117	159	180	176

Digital images ?

- For an image I
 - Index (0,0): Top left corner
 - $I(x,y)$: intensity of pixel at the position (x,y)

x =	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72
y =	41	210 209 204 202 197 247 143 71 64 80 84 54 54 57 58													
	42	206 196 203 197 195 210 207 56 63 58 53 53 61 62 51													
	43	201 207 192 201 198 213 156 69 65 57 55 52 53 60 50													
	44	216 206 211 193 202 207 208 57 69 60 55 77 49 62 61													
	45	221 206 211 194 196 197 220 56 63 60 55 46 97 58 106													
	46	209 214 224 199 194 193 204 173 64 60 59 51 62 56 48													
	47	204 212 213 208 191 190 191 214 60 62 66 76 51 49 56													
	48	214 215 207 208 180 172 188 69 72 55 49 56 52 56													
	49	209 205 214 205 204 196 187 196 86 62 66 87 57 60 48													
	50	208 209 205 203 202 186 174 185 149 71 63 55 55 45 56													
	51	207 210 211 199 217 194 183 177 209 90 62 64 52 93 52													
	52	208 205 209 209 197 194 183 187 187 239 58 68 61 51 56													
	53	204 206 203 209 195 203 188 185 183 221 75 61 58 60 60													
	54	200 203 199 236 188 197 183 190 183 196 122 63 58 64 66													
	55	205 210 202 203 199 197 196 181 173 186 105 62 57 64 63													



Digital images ?

- Principal type of images

- Binary image:

- $I(x,y) \in \{0, 1\}$
 - 1 pixel: 1 bit



- Gray image:

- $I(x,y) \in [0..255]$
 - 1 pixel: 8 bits (1 byte)



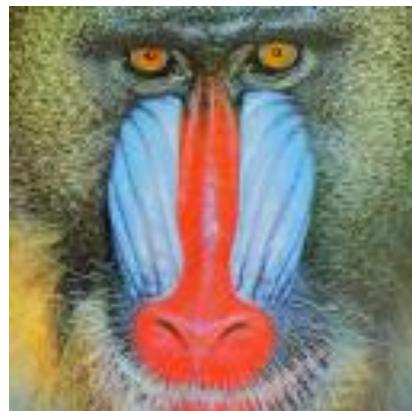
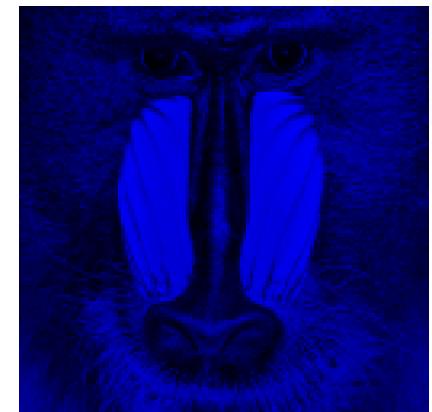
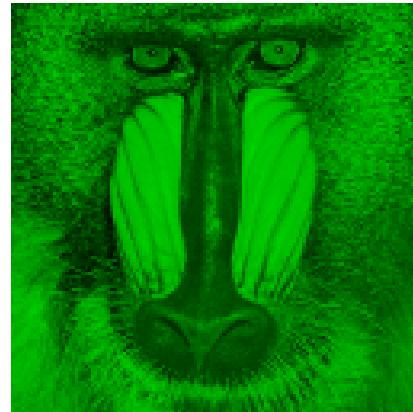
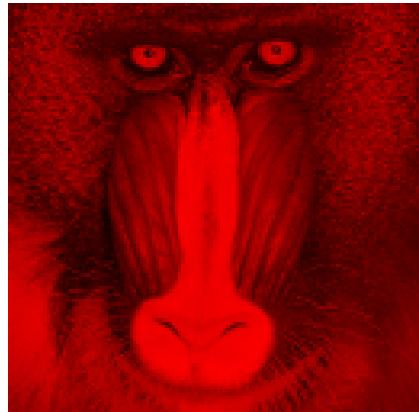
- Color image

- $I_R(x,y), I_G(x,y), I_B(x,y) \in [0..255]$
 - 1 pixel: 24 bits (3 bytes)



- Other : multi-spectre, depth image,...

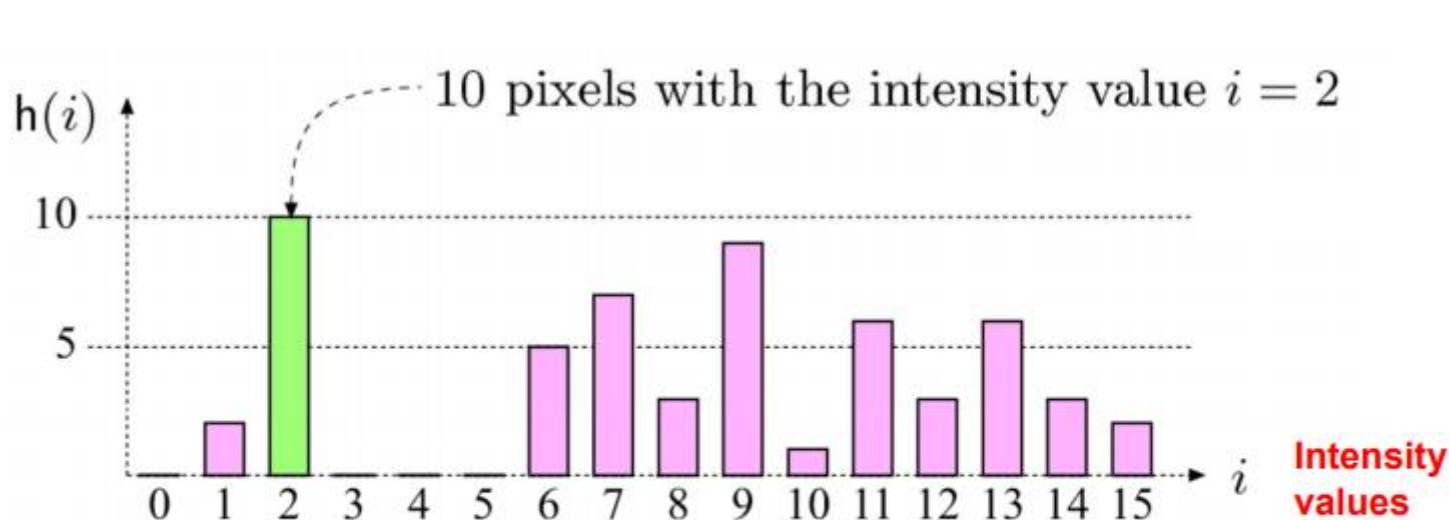
Color image in RGB space



It exists other color spaces:
Lab, HSV, ...

Image histogram

- Histogram is a graphical representation of the repartition of colours among the pixels of a numeric image.



$h(i)$	0	2	10	0	0	0	5	7	3	9	1	6	3	6	3	2
i	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Image histogram

- Histogram
 - Should be normalized by dividing all elements to total number of pixels in the image

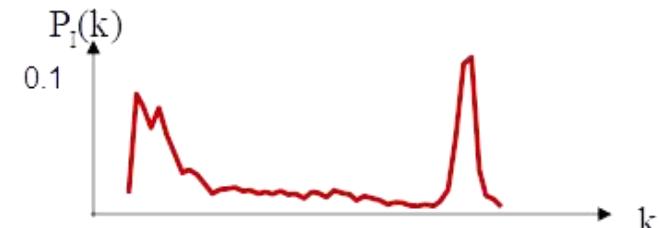
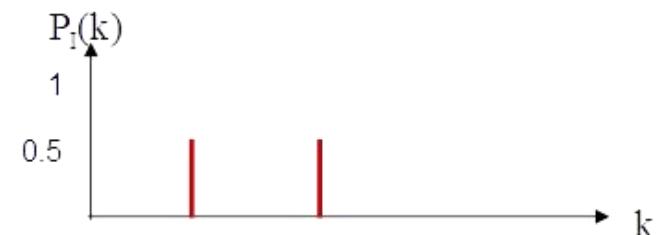


Image dynamic range = [min_value, max_value]

Image histogram

- Histogram
 - Only statistic information
 - No indication about the location of pixel (no spatial information)
 - Different image can have the same histogram

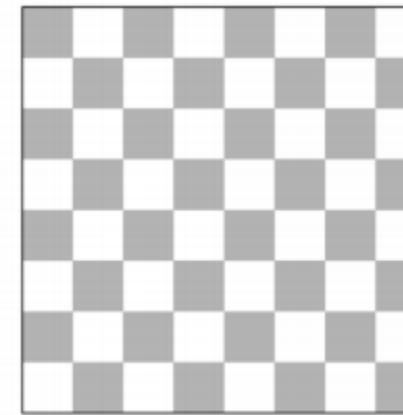
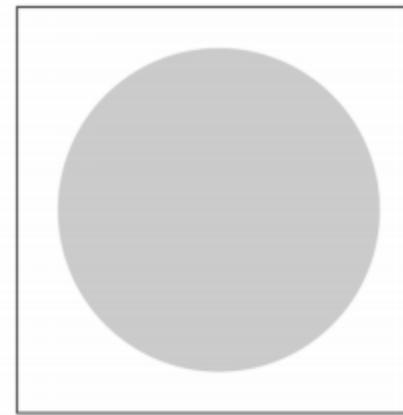
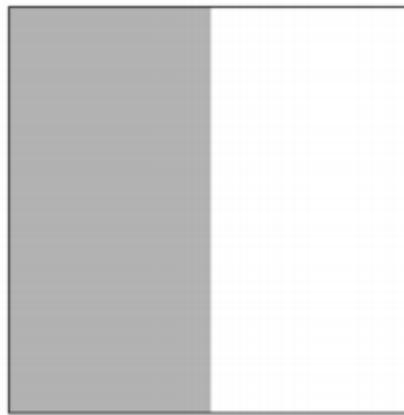


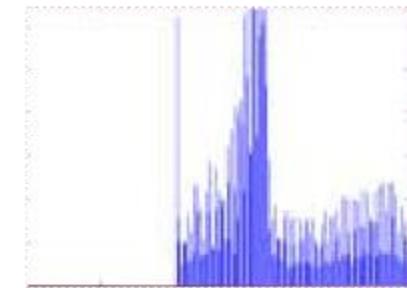
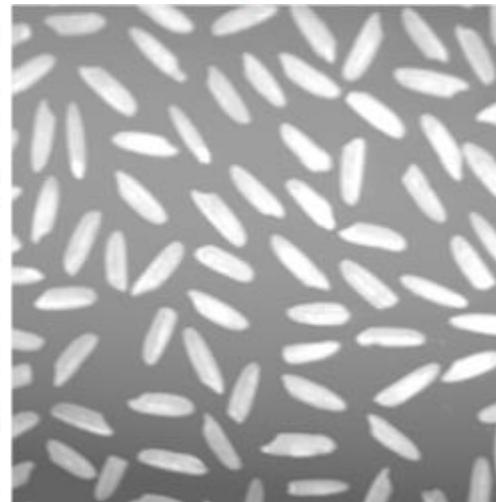
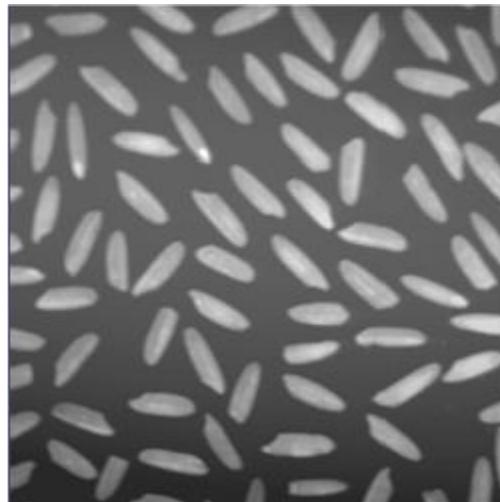
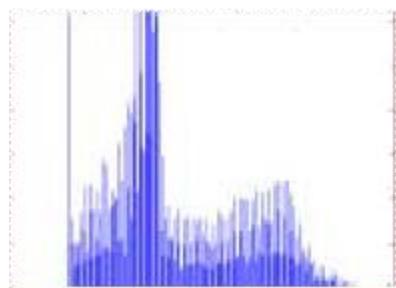
Image Brightness

- Brightness of a grayscale image is the average intensity of all pixels in an image
 - refers to the overall lightness or darkness of the image

$$B(I) = \frac{1}{wh} \sum_{v=1}^h \sum_{u=1}^w I(u, v)$$

Divide by total number of pixels

Sum up all pixel intensities



Contrast

- The contrast of a grayscale image indicate how easily object in the image can be distinguished
- Many different equations for contrast exist
 - Standard deviation of intensity values of pixels in the image

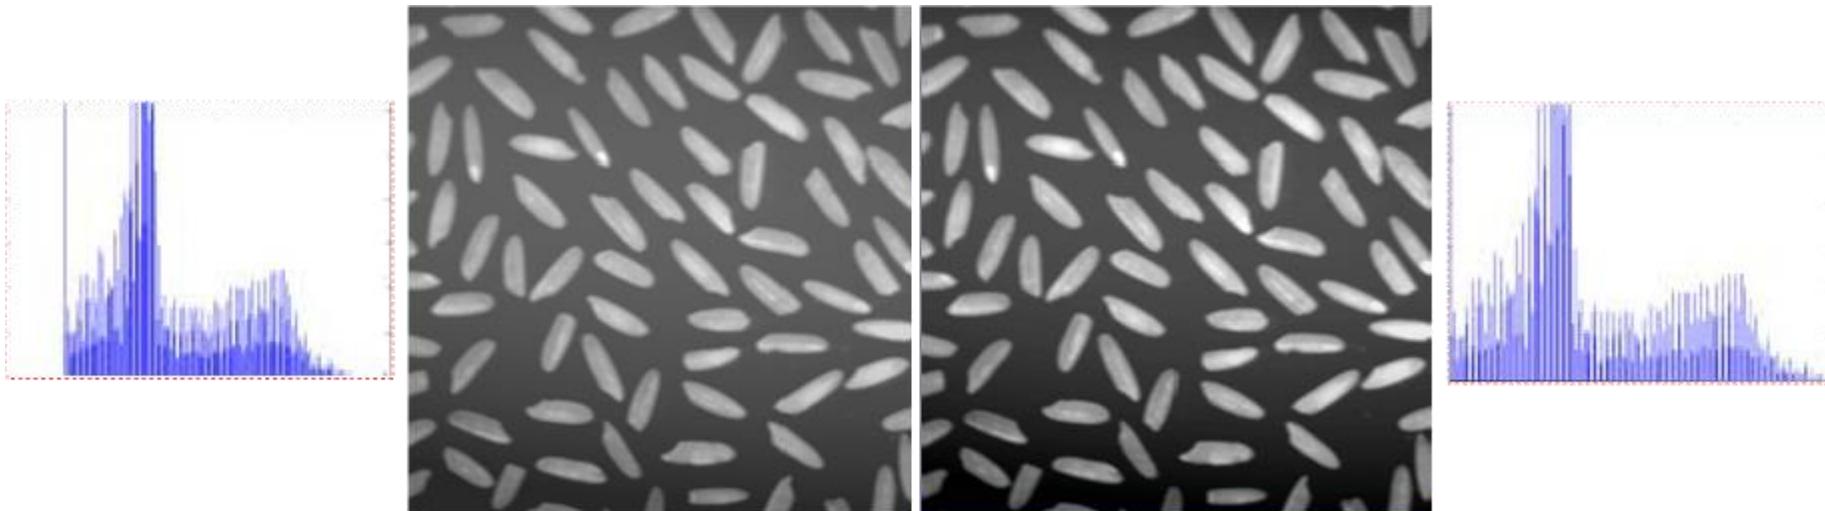
$$C = \sqrt{\frac{1}{M \times N} \sum_{x=0}^{N-1} \sum_{y=0}^{M-1} (f(x,y) - Moy)^2}$$

- Difference between intensity value maximum et minimum

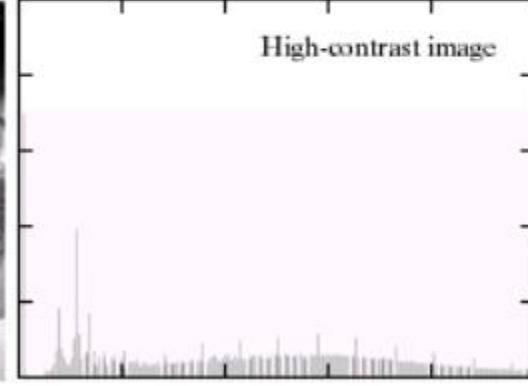
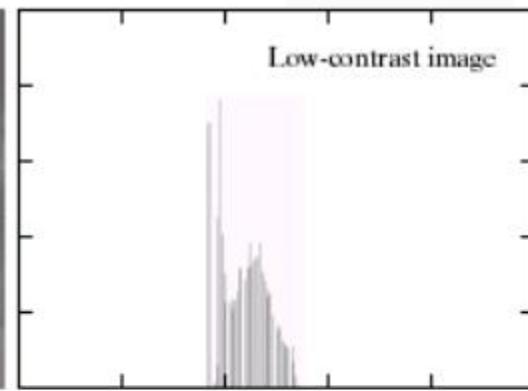
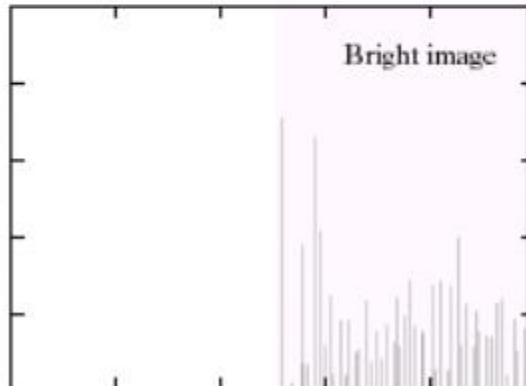
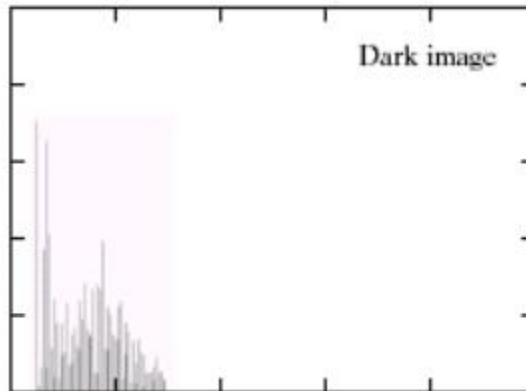
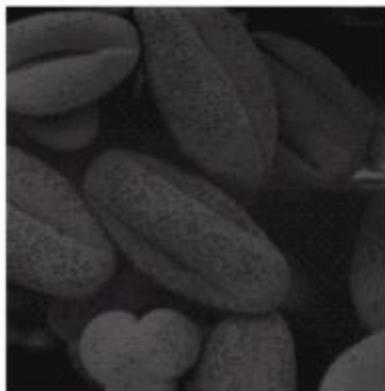
$$C = \frac{\max[f(x,y)] - \min[f(x,y)]}{\max[f(x,y)] + \min[f(x,y)]}$$

Contrast

- Contrast vs histogram



Examples



Contrast Enhancement

- Modify pixel intensities to obtain higher contrast
- There are several methods:
 - Linear stretching of intensity range:
 - Linear transform
 - Linear transform with saturation
 - Piecewise linear transform
 - Non-linear transform (Gama correction)
 - Histogram equalization (Cân bằng histogram)

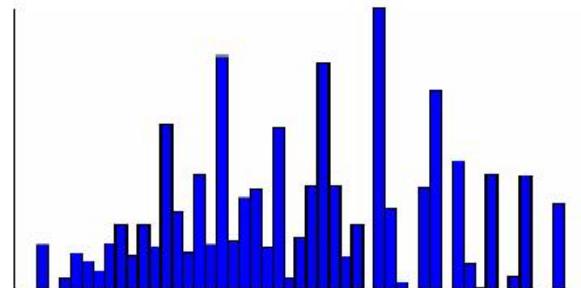
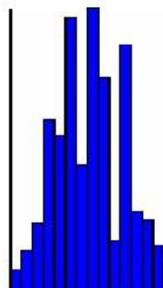
Linear stretching



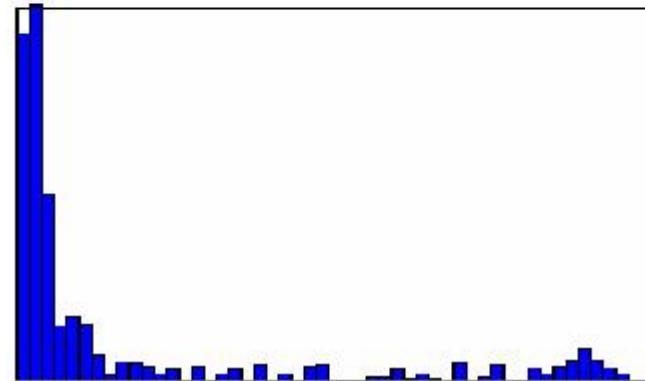
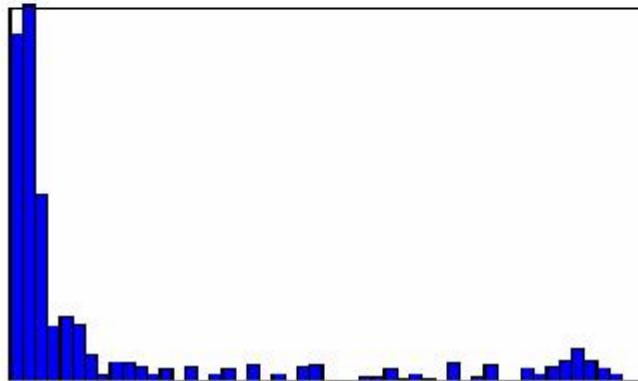
$$f_{ac}(a) = a_{min} + (a - a_{low}) \cdot \frac{a_{max} - a_{min}}{a_{high} - a_{low}}$$

If $a_{min} = 0$ and $a_{max} = 255$

$$f_{ac}(a) = (a - a_{low}) \cdot \frac{255}{a_{high} - a_{low}}$$



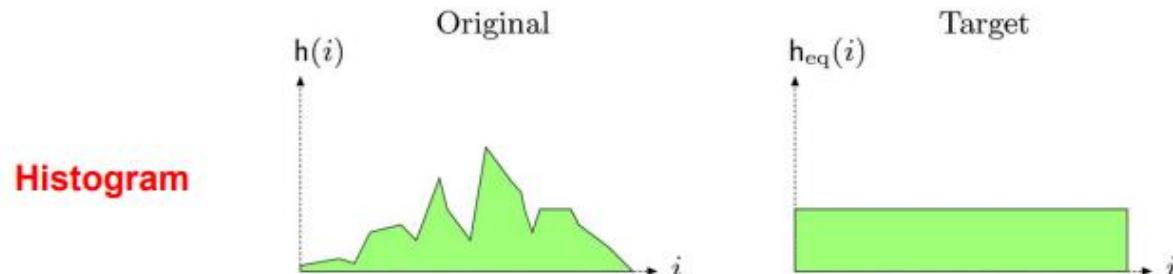
Linear stretching



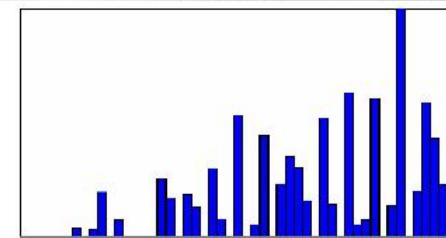
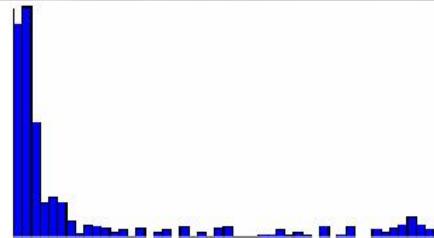
No efficace?

Histogram equalization

- Change histogram of modified image into uniform distribution



- No parameters. OpenCV: `cv2.equalizeHist(img)`



Color Image histogram

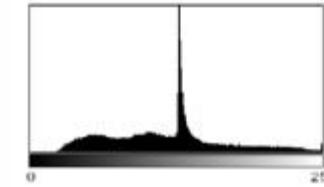
- Intensity histogram:

- Convert color image to grayscale

=> Compute histogram of gray scale image



(a)

(b) h_{Lum} 

(c) R

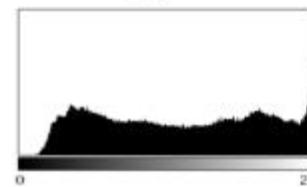
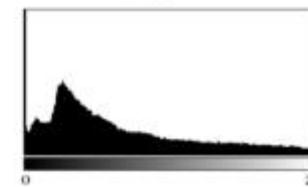
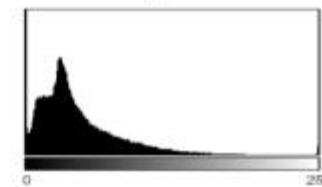


(d) G



(e) B

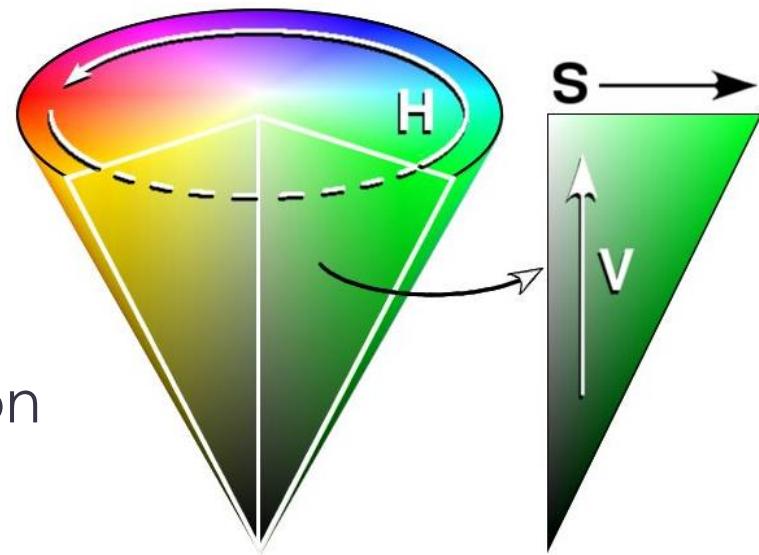
3 histograms for (R,G,B)

(f) h_R (g) h_G (h) h_B

- 3D histogram: a color identified by 3 values. Not usually because of big elements

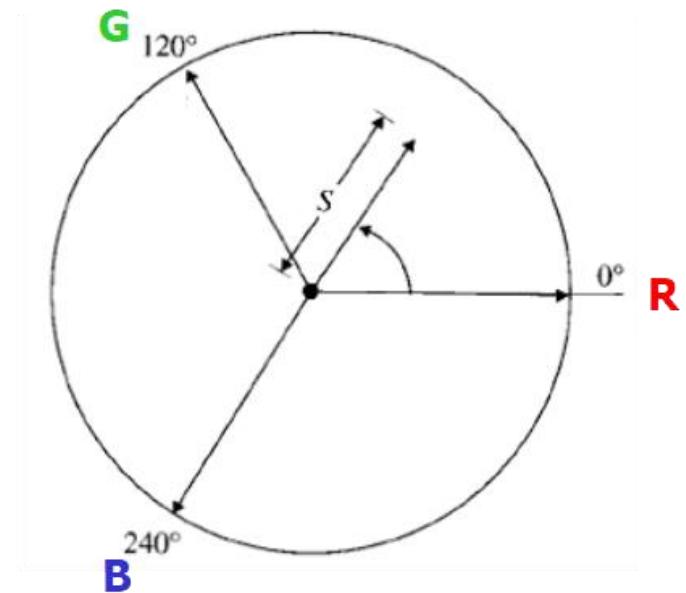
HSV (Hue – Saturation- Value)

- The Hue-Saturation-Value (HSV) color space is useful for segmentation and recognition
 - Non-linear conversion
 - Visual representation of colors
- We identify for a pixel:
 - The pixel *intensity (value)*
 - The pixel *color (hue + saturation)*
- RGB does not have this separation



HSV (Hue – Saturation- Value)

- **Hue (H)** is coded as an angle between 0 and 360
- **Saturation (S)** is coded as a radius between 0 and 1
 - $S = 0$: gray
 - $S = 1$: pure color
- **Value (V) = MAX (Red, Green, Blue)**



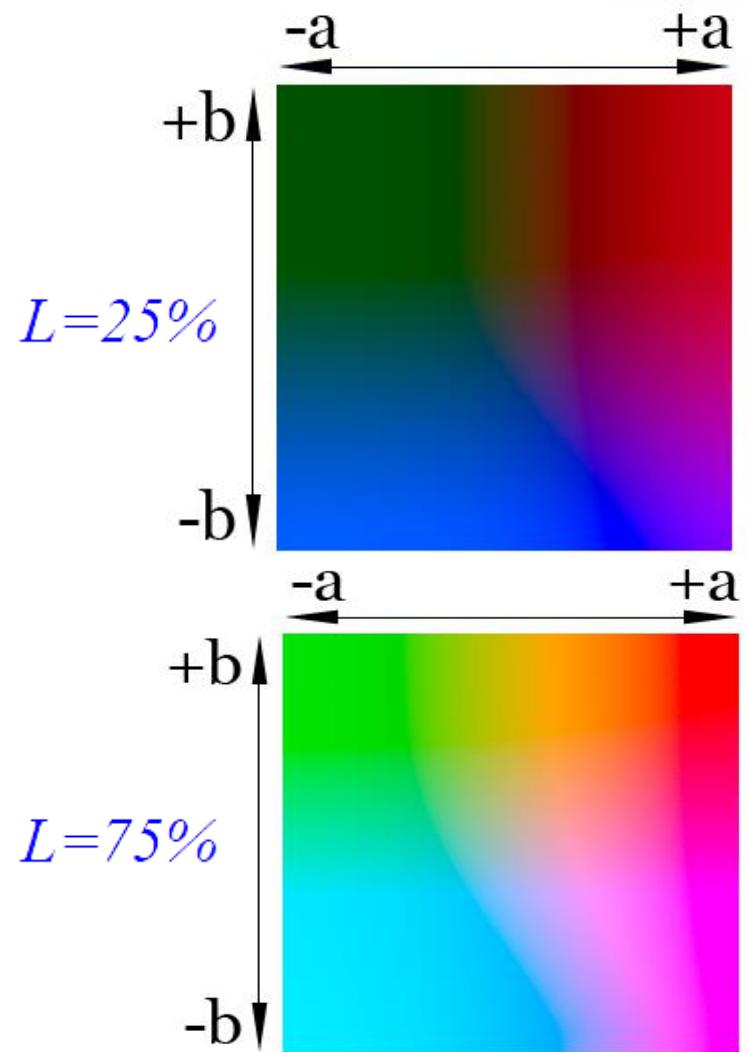
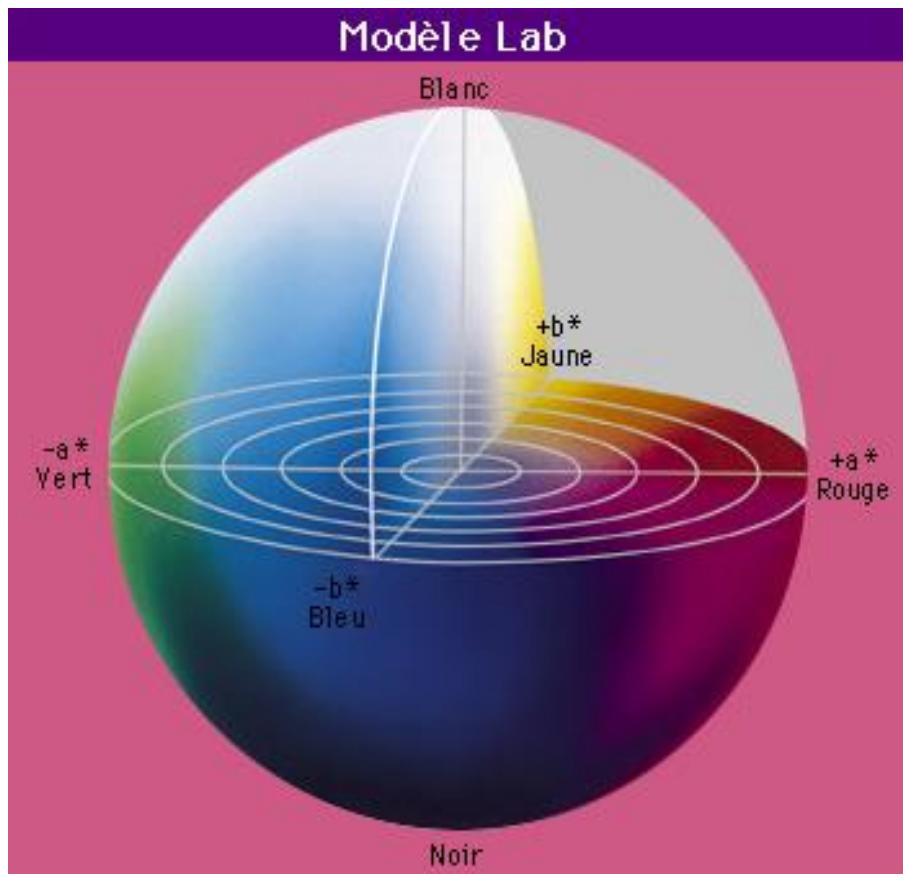
HSV (Hue – Saturation- Value)

- If we know the color of the object we are looking for, we can model it using a **hue interval**
- Take care, because it is an angle (periodic value)
 - Hue < 60° means nothing
 - Is 350° smaller or bigger than 60°?
 - Define an interval: $350^\circ < \text{Hue} < 60^\circ$ (for example)
- This interval is valid if Saturation > threshold (otherwise gray level)
- This is independant of **Value** , which is more sensible to light conditions

Lab color space

- The **Lab** system (sometimes **L*a*b***) is based on a study from human vision
 - independant from all technologies
 - presenting colors as seen by the human eyes
- Colors are defined using 3 values
 - L is the luminance, going from 0% (black) to 100% (white)
 - a* represents an axis going from green (negative value, -127) to red (positive value, +127)
 - b* represents an axis going from blue (negative value, -127) to yellow (positive value, +127)

Lab color space

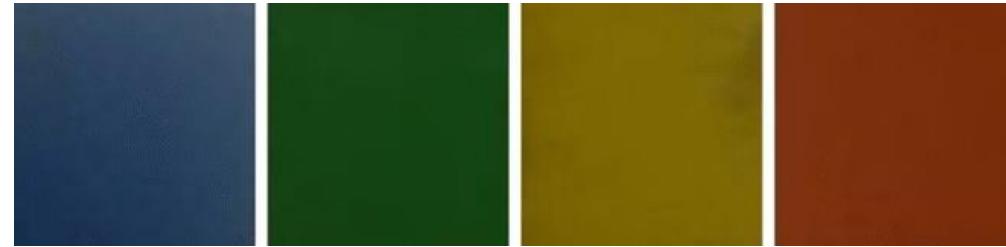


Color space vs illumination conditions

- collected 10 images of the cube under varying illumination conditions



- separately cropped every color to get 6 datasets for the 6 different colors



Changes in color due to varying Illumination conditions

- Compute the density plot: Check the distribution of a particular color say, blue or yellow in different color spaces. The density plot or the 2D Histogram gives an idea about the variations in values for a given color

Color space vs illumination conditions

- Similar illumination: very compact

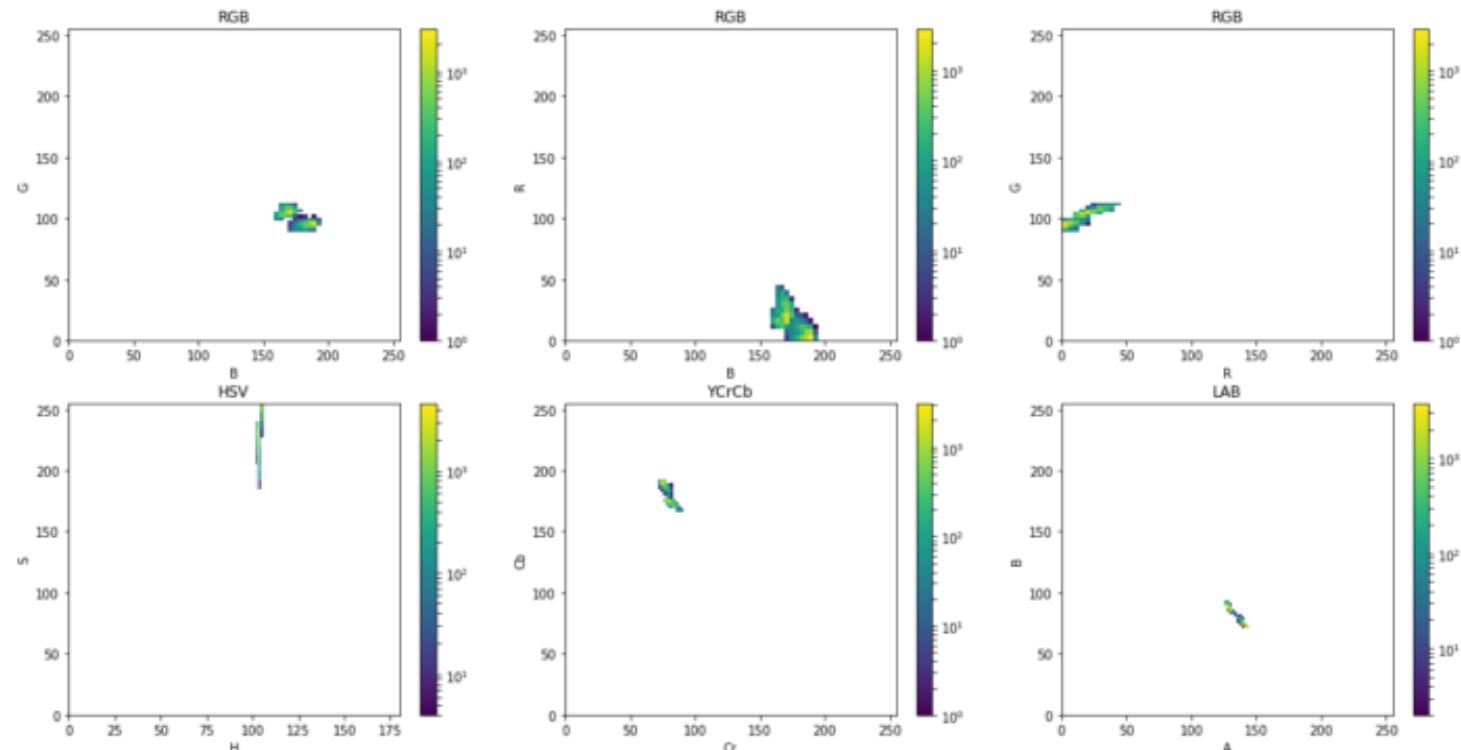


Fig.: Density Plot showing the variation of values in color channels for 2 similar bright images of **blue color**

Color space vs illumination conditions

- Similar illumination: very compact

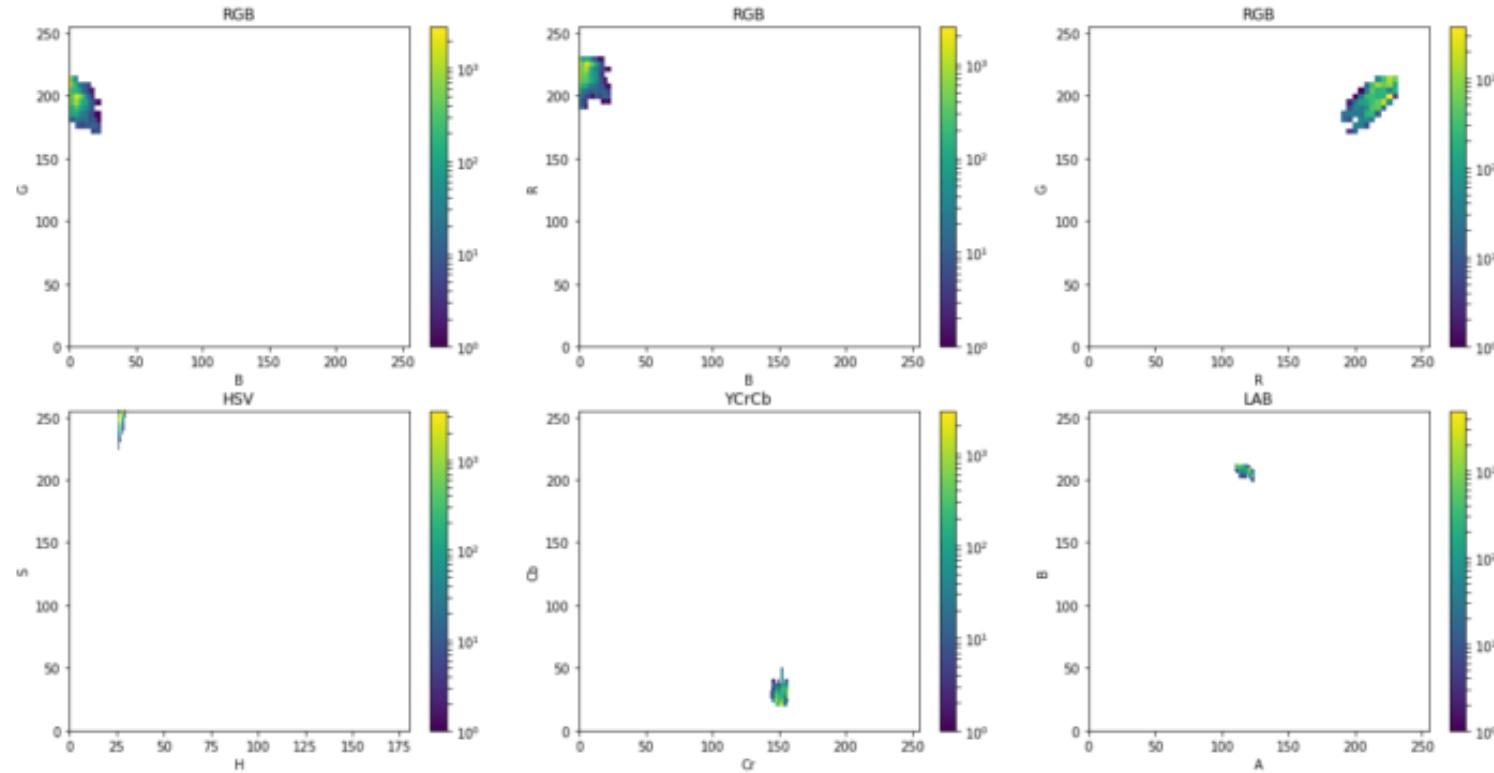


Fig.: Density Plot showing the variation of values in color channels for 2 similar bright images of **yellow color**

Color space vs illumination conditions

- Different illumination:

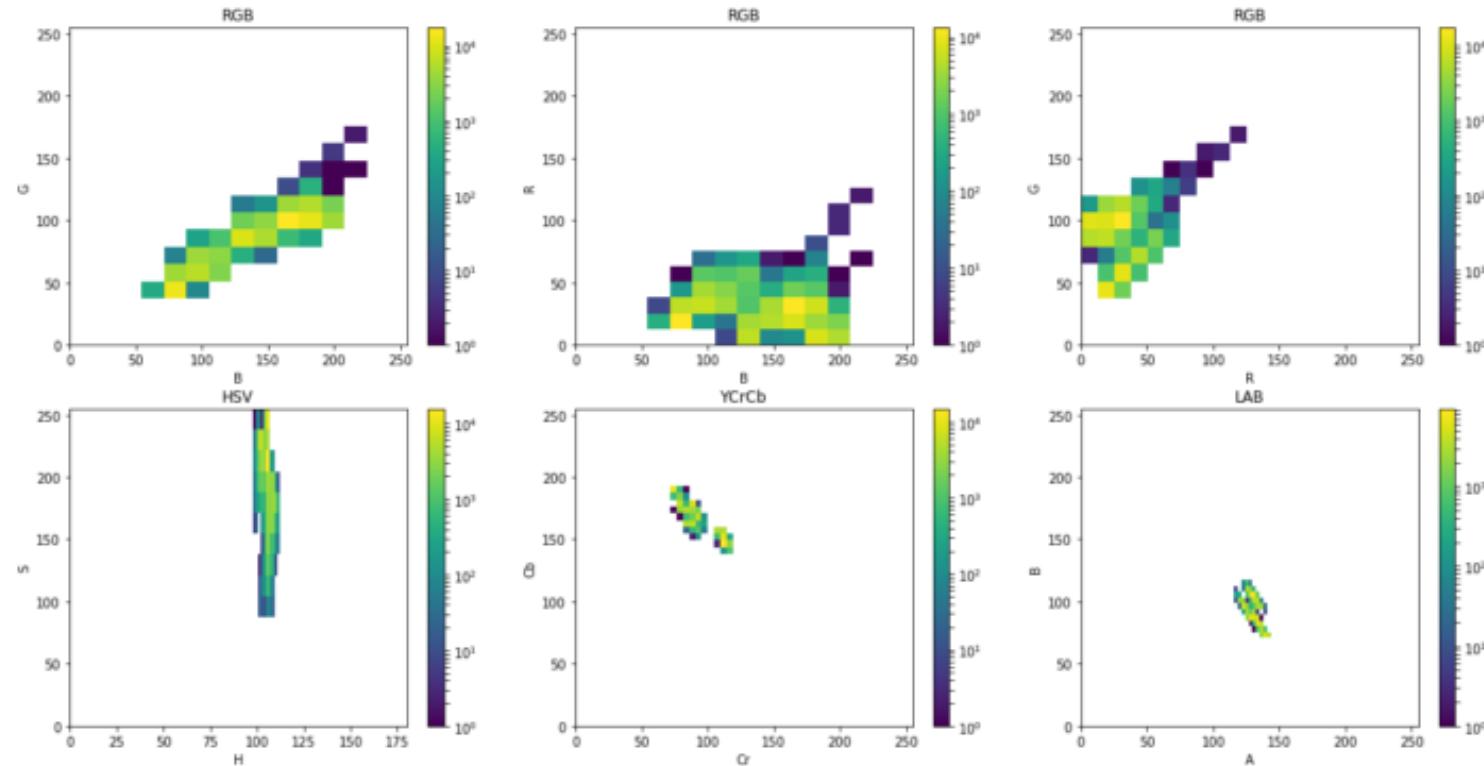


Fig.: Density Plot showing the variation of values in color channels under varying illumination for the **blue color**

Color space vs illumination conditions

- Different illumination:

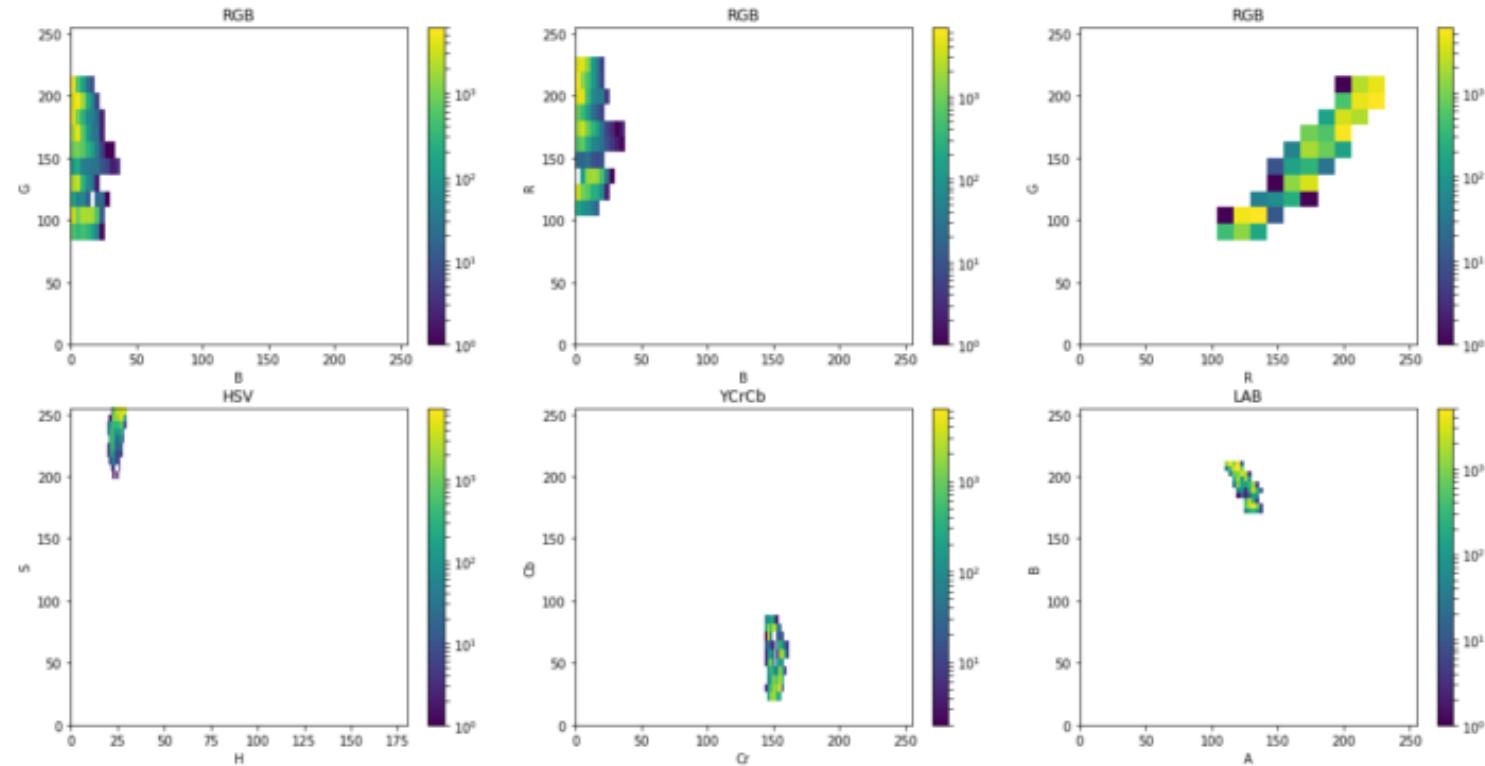


Fig.: Density Plot showing the variation of values in color channels under varying illumination for the **yellow color**

Color space vs illumination conditions

- Different illumination:
 - RGB space: the variation in the value of channels is very hight
 - HSV: compact in **H**. Only H contains information about the absolute color → a choix
 - YCrCb, LAB: compact in **CrCb** and in **AB**
 - Higher level of compactness is in LAB
 - Convert to other color spaces (OpenCV):
 - cvtColor(bgr, ycb, COLOR_BGR2YCrCb);
 - cvtColor(bgr, hsv, COLOR_BGR2HSV);
 - cvtColor(bgr, lab, COLOR_BGR2Lab);

Spatial convolution

- Image filtering : For each pixel, compute function of local neighborhood and output a new value
 - **Same function** applied at each position
 - Output and input image are typically **the same size**
- Convolution : Linear filtering, function is a weighted sum/difference of pixel values

$$I' = I * K$$

- Really important!
 - Enhance images: Denoise, smooth, increase contrast, etc.
 - Extract information from images:
 - Texture, edges, distinctive points, etc.
 - Detect patterns
 - Template matching

Spatial convolution



Original image

$$\text{Original image} * \text{Mask (kernel)} = \text{Filtered image}$$

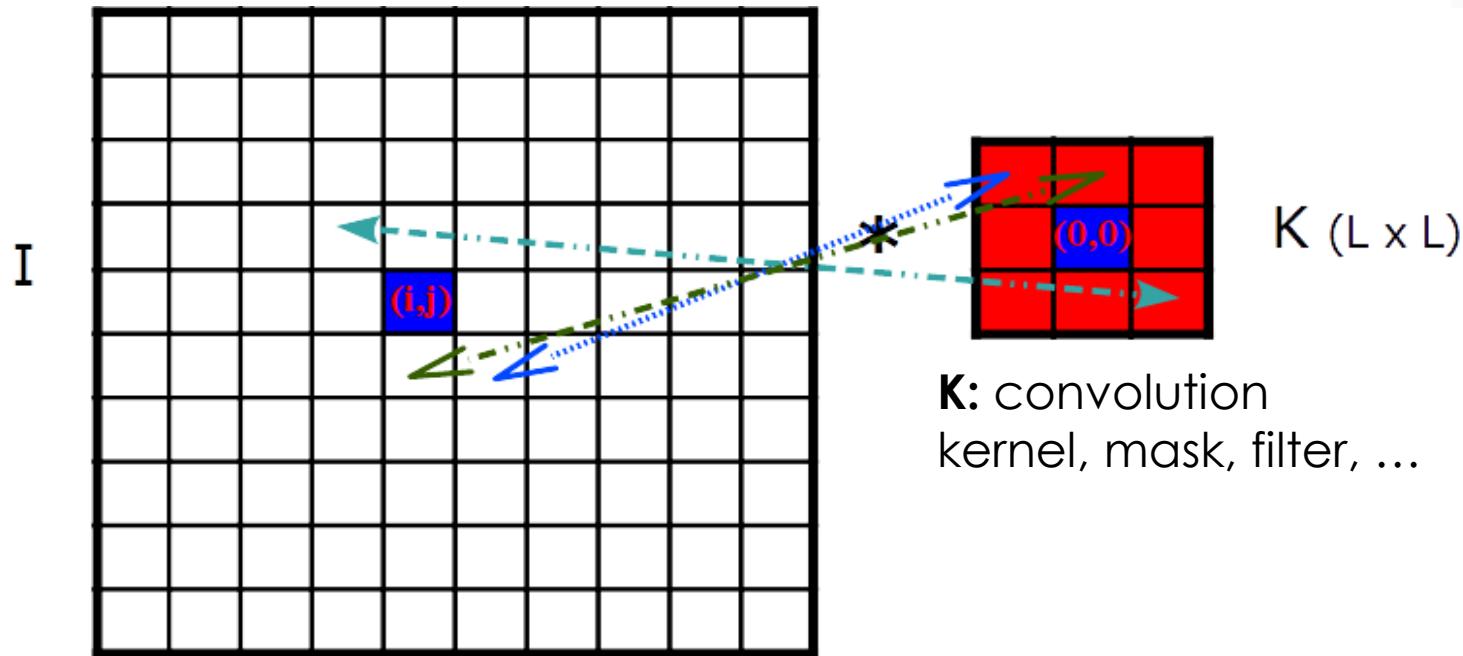
Mask (kernel)



Filtered image

Spatial convolution

- New value of a pixel (i,j) is a weighted sum of its neighbors



$$I'(i, j) = \sum_{u=-\frac{(L-1)}{2}}^{\frac{(L-1)}{2}} \sum_{v=-\frac{(L-1)}{2}}^{\frac{(L-1)}{2}} I(i-u, j-v) K(u, v)$$

Spatial convolution

- New value of a pixel(i,j) is a weighted sum of its neighbors

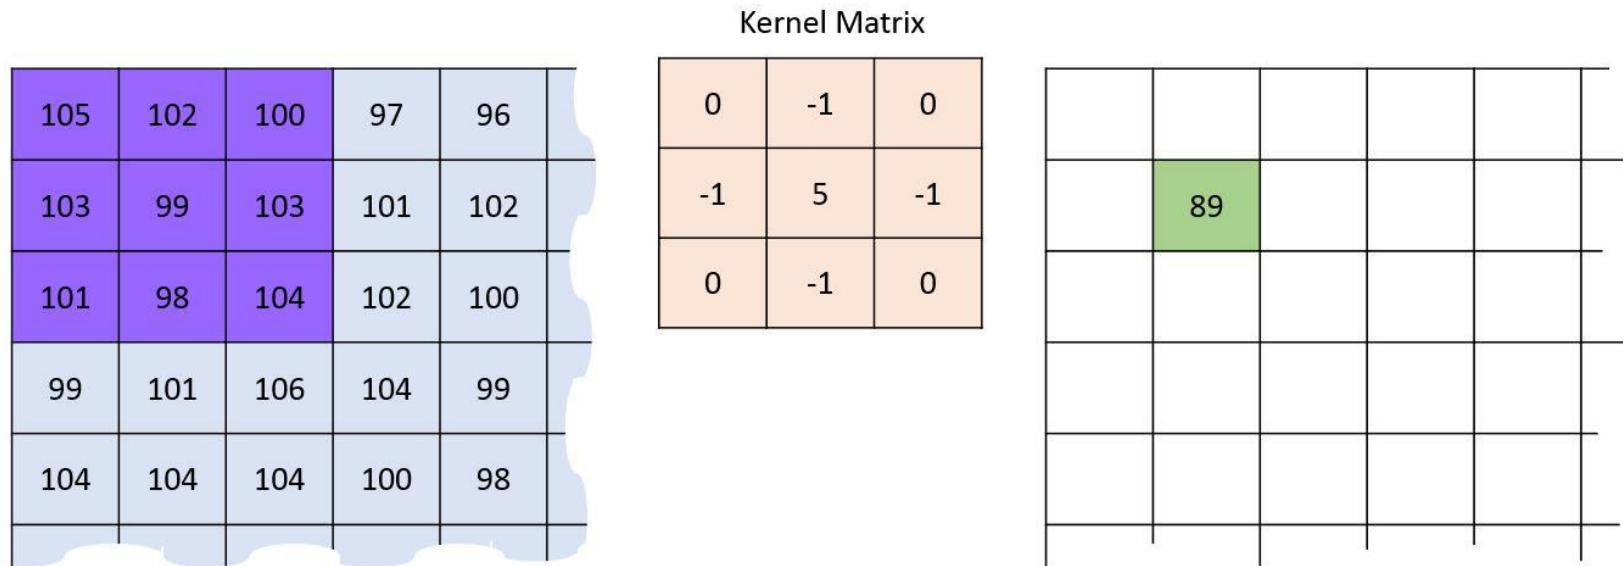


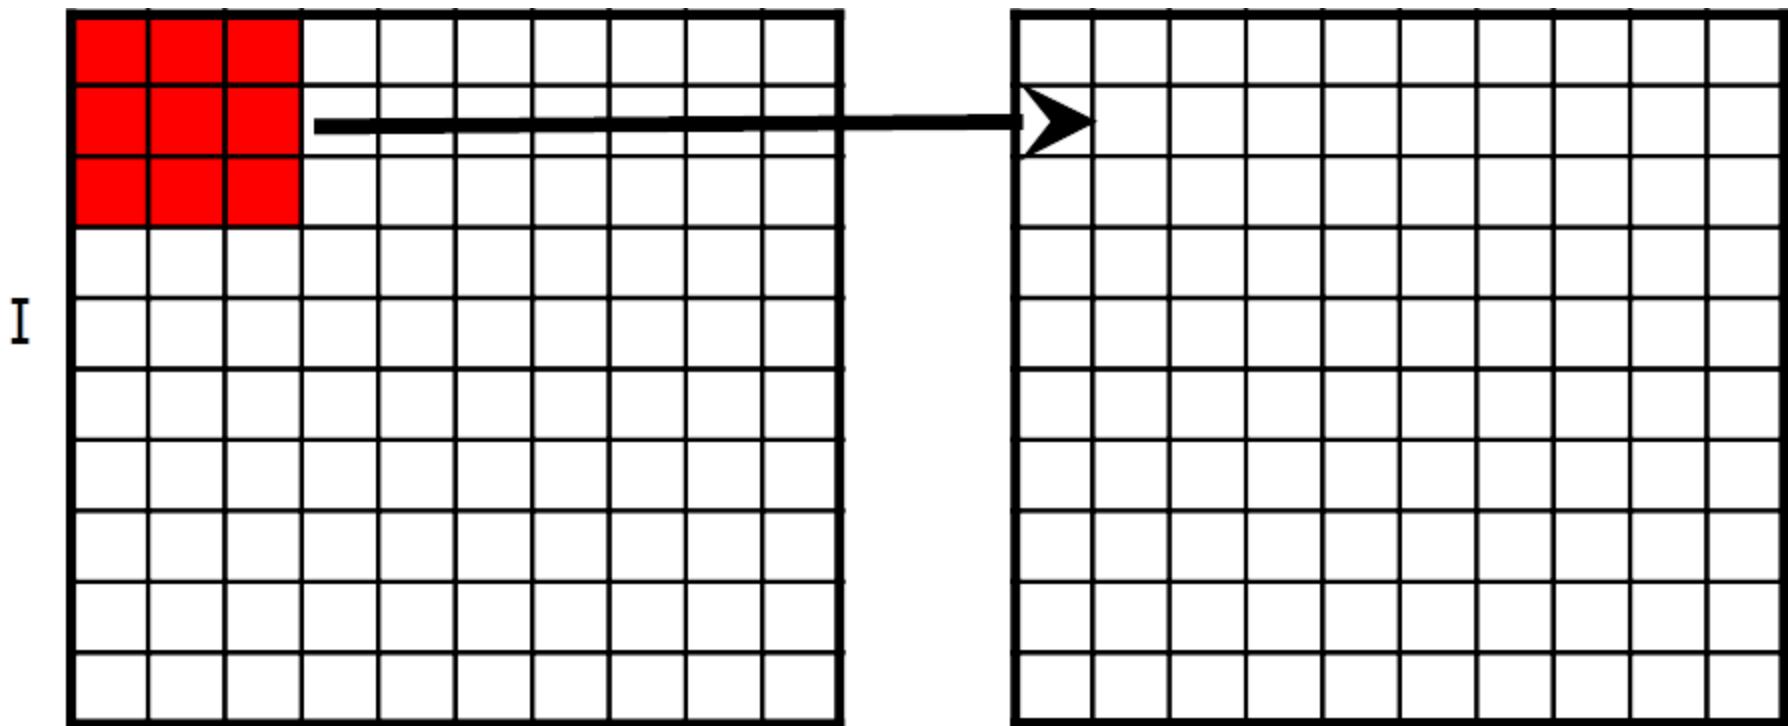
Image Matrix

$$\begin{aligned}
 & 105 * 0 + 102 * -1 + 100 * 0 \\
 & + 103 * -1 + 99 * 5 + 103 * -1 \\
 & + 101 * 0 + 98 * -1 + 104 * 0 = 89
 \end{aligned}$$

Output Matrix

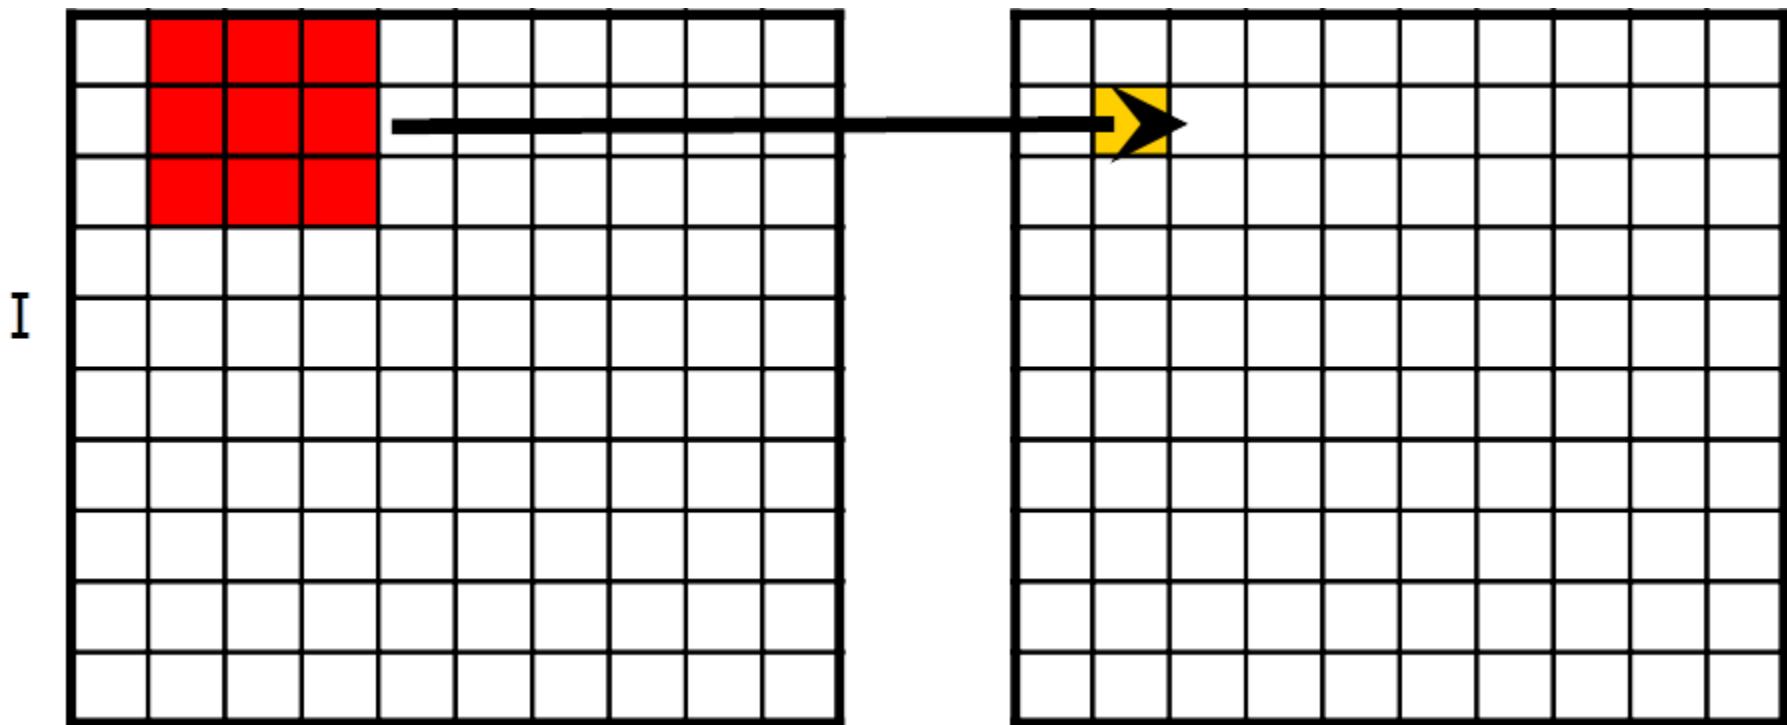
Spatial convolution

$$I' = I * K$$



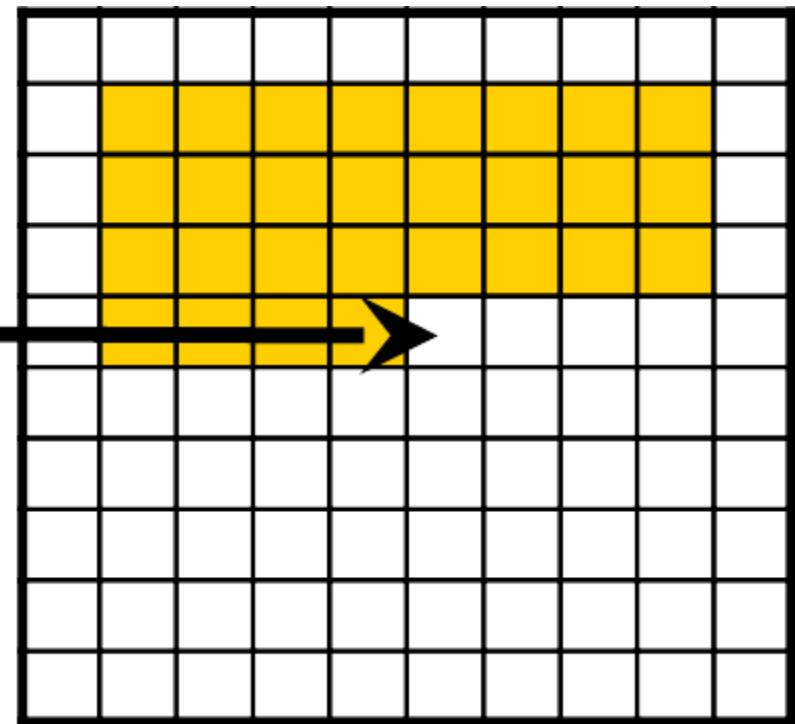
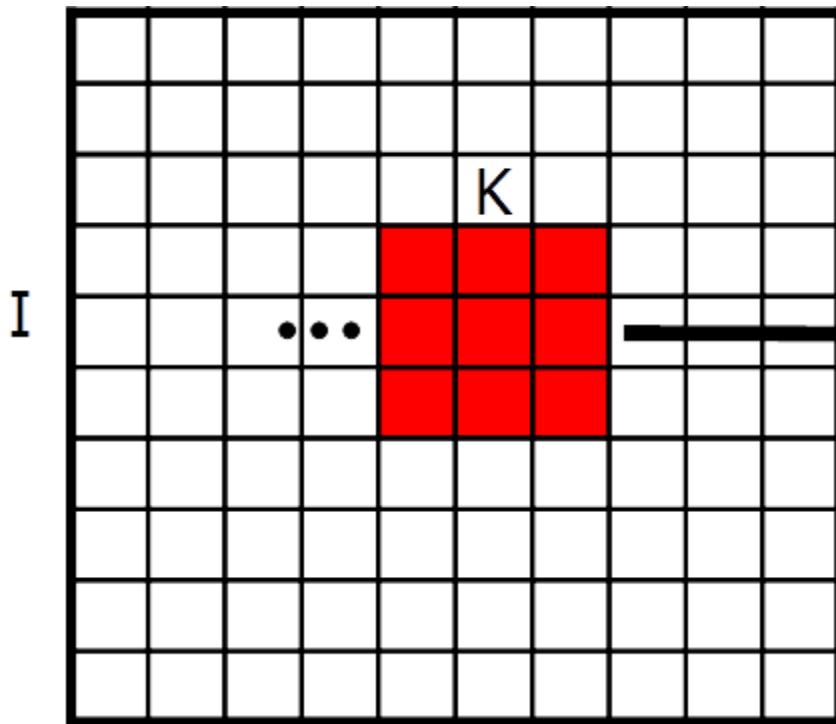
Spatial convolution

$$I' = I * K$$



Spatial convolution

$$I' = I * K$$



Spatial convolution

- Border problem?
 - Zero padding in the input matrix
 - reflect across edge:
 - $f(-x,y) = f(x,y)$
 - $f(-x,-y) = f(x,y)$
 - ...

?	?	?	?	?	?	?	?	?	?	?
?										?
?										?
?										?
?										?
?										?
?										?
?	?	?	?	?	?	?	?	?	?	?

0	0	0	0	0	0	
0	105	102	100	97	96	
0	103	99	103	101	102	
0	101	98	104	102	100	
0	99	101	106	104	99	
0	104	104	104	100	98	

105	105	102	100	97	96	
105	105	102	100	97	96	
103	103	99	103	101	102	
101	101	98	104	102	100	
99	99	101	106	104	99	
104	104	104	104	100	98	



Spatial convolution

0	0	0	0	0	0	
0	105	102	100	97	96	
0	103	99	103	101	102	
0	101	98	104	102	100	
0	99	101	106	104	99	
0	104	104	104	100	98	

Kernel Matrix

0	-1	0
-1	5	-1
0	-1	0

320					
210	89	111			

Image Matrix

$$\begin{aligned}
 & 0 * 0 + 0 * -1 + 0 * 0 \\
 & + 0 * -1 + 105 * 5 + 102 * -1 \\
 & + 0 * 0 + 103 * -1 + 99 * 0 = 320
 \end{aligned}$$

Output Matrix

105	105	102	100	97	96	
105	105	102	100	97	96	
103	103	99	103	101	102	
101	101	98	104	102	100	
99	99	101	106	104	99	
104	104	104	104	100	98	

Kernel Matrix

0	-1	0
-1	5	-1
0	-1	0

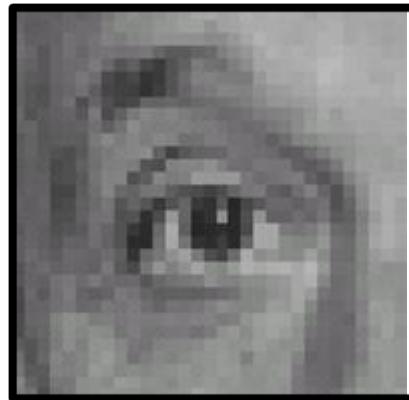
110					
	89	111			

Image Matrix

Some kernels

- 2D spatial convolution
 - is mostly used in image processing for feature extraction
 - And is also the core block of Convolutional Neural Networks (CNNs)
- Each kernel has its own effect and is useful for a specific task such as
 - blurring (noise removing),
 - sharpening,
 - edge detection,
 -

Some kernels



Original image

*

0	0	0
0	1	0
0	0	0



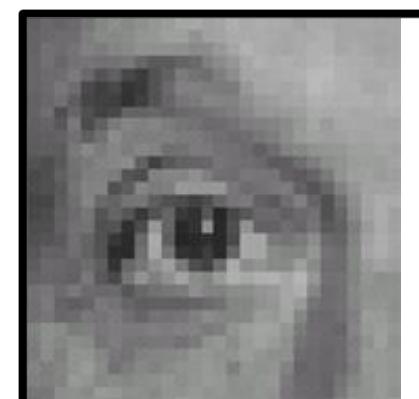
Filtered image
(no change)



Original image

*

0	0	0
1	0	0
0	0	0



Filtered image
(shifted left by 1 pixel)

Some kernels

- Box filter (**mean filter**):
 - Replace each pixel with an average of its neighborhood
 - Achieve smoothing effect

$$\frac{1}{9} \times \begin{array}{|c|c|c|}\hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline\end{array}$$



Original image



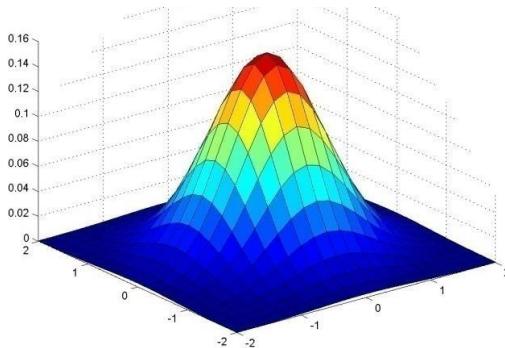
Filtered image
with box size 5x5



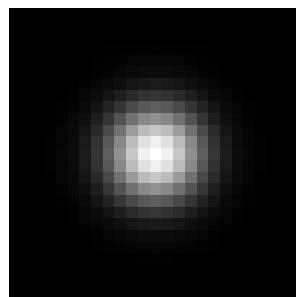
Filtered image
with box size 11x11

Some kernels

- Gaussian filter



Gaussian function in 3D



Gaussian image

0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

Gaussian filter with size 5 x 5 ,
sigma = 1

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

Rule for Gaussian filter:
set **filter half-width to about 3σ**

Some kernels

- Gaussian filter:
 - **Low-pass filter:** Remove high-frequency component from the image
 - Image becomes more smooth
 - Better than mean filter
 - Convolution with itself is an other Gaussian
 - Repeat the conv. With small-width kernel => get the same result as larger-width kernel would have.
 - Convolving **2 times** with Gaussian kernel of **width σ** is the same as convolving **once with kernel of width $\sigma\sqrt{2}$** : $I * G_\sigma * G_\sigma = I * G_{\sigma\sqrt{2}}$
 - **Separable filter:** The 2D Gaussian can be expressed as the product of 2 functions : one function of x and a function of y:
 - $G_\sigma(x,y) = G_\sigma(x).G_\sigma(y)$

Some kernels

- Gaussian filter



Original image



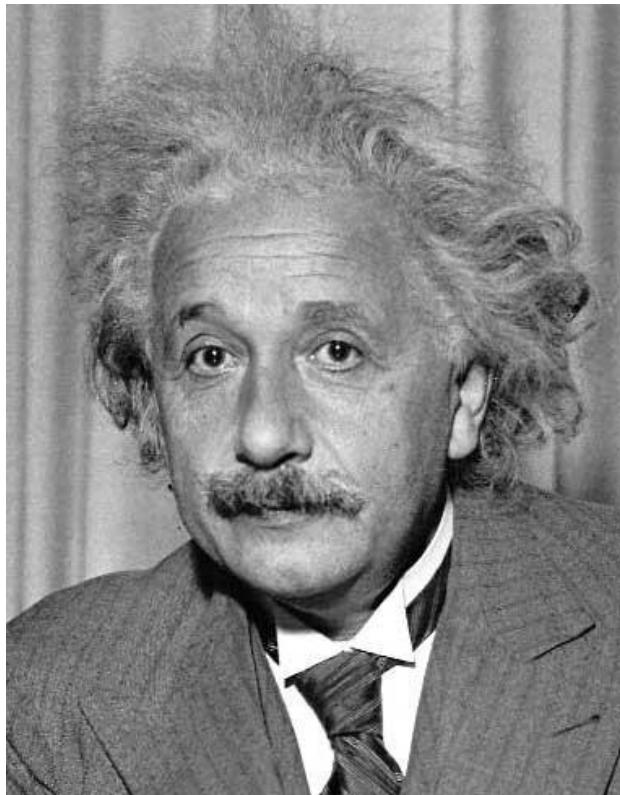
Filtered image
with box size 5x5



Filtered image
with box size 11x11

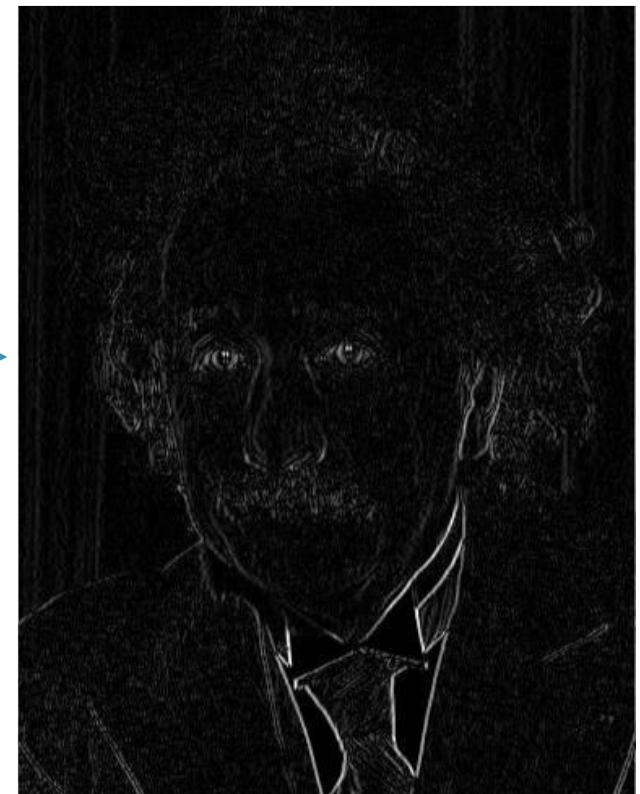
Some kernels

■ Sobel



*

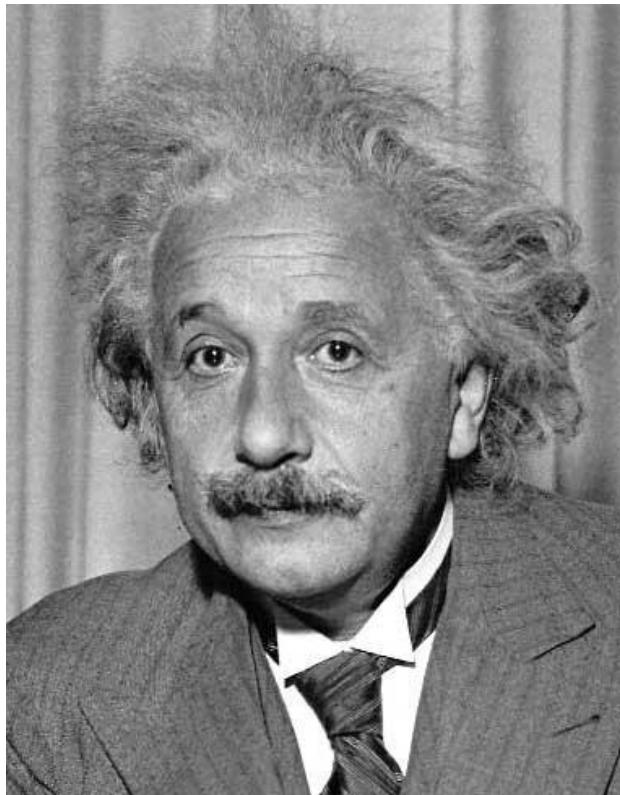
-1	0	1
-2	0	2
-1	0	1



Vertical Edge
(absolute value)

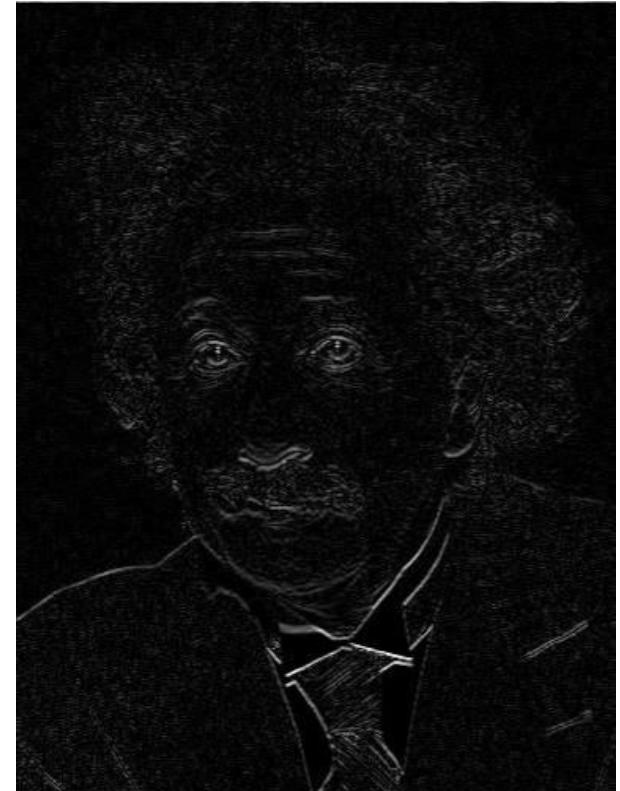
Some kernels

- Sobel



*

-1	-2	-1
0	0	0
1	2	1

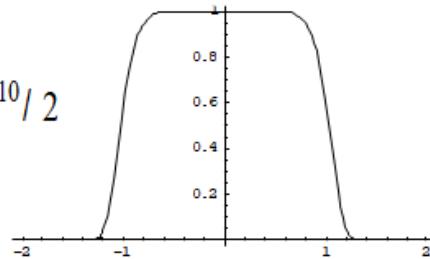


Horizontal Edge
(absolute value)

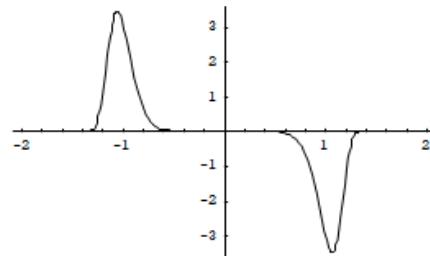
Edge detection

- Edges are corresponding to:
 - Maximums of the first derivative
 - Zero-crossing in the second derivative

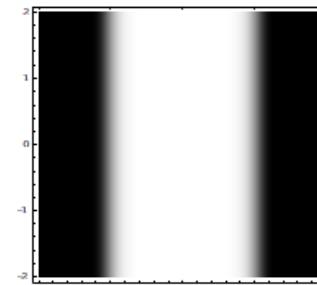
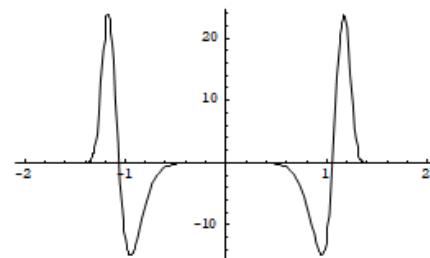
$$f(x, y) = e^{-x^{10}/2}$$



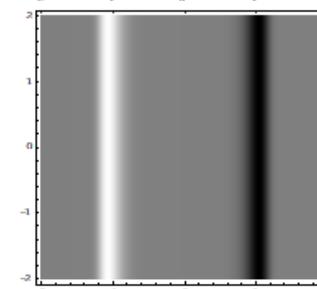
$$\frac{\partial f}{\partial x}$$



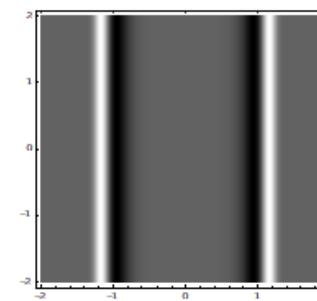
$$\frac{\partial^2 f}{\partial x^2}$$



Image



First derivative



Second derivative

Edge detection

- Use the first derivatives of the image:
 - Compute the first derivatives (with different kernels: **Sobel**, **Prewitt**, **Robert**)
 - Find local maximum
 - Edge composed of pixels having **maximum value of the first derivatives** of image
 - Can **use a threshold** to detect edge rapidly
 - Can make several steps to obtain the optimal edge: **Canny detector**
- Use second derivative
 - Compute the second derivative (**Laplacian** filter)
 - Find zero-crossing

Edge detection with first derivatives

- Filters used to compute the first derivatives of image

- Robert

1	0
0	-1

0	1
-1	0

- Prewitt

- less sensitive to noise
- Smoothing with mean filter,
then compute 1st derivative

-1	-1	-1
0	0	0
1	1	1

-1	0	1
-1	0	1
-1	0	1

- Sobel:

- less sensitive to noise
- Smoothing with gaussian,
then computing 1st derivative

-1	-2	-1
0	0	0
1	2	1

y

-1	0	1
-2	0	2
-1	0	1

x

Edge detection with first derivatives

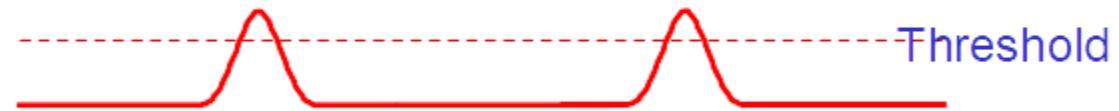
Image 1D $f(x)$



1st derivative $f'(x)$



$|f'(x)|$



Edge pixels:

$|f'(x)| > \text{Threshold}$



Image derivatives

- 1st derivatives :

$I *$

-1	0	1
-2	0	2
-1	0	1



I_x

First derivative
of image with
respect to x

$I *$

-1	-2	-1
0	0	0
1	2	1



I_y

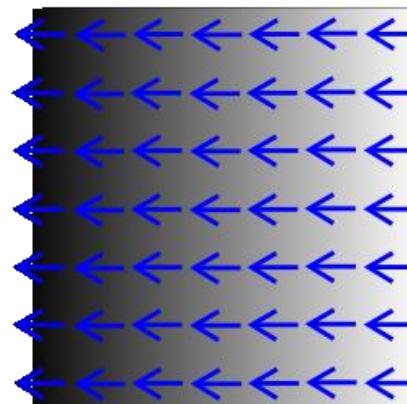
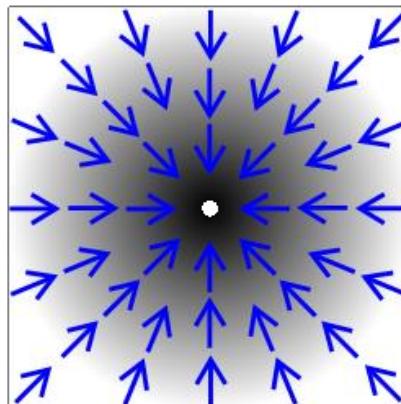
First derivative
of image with
respect to y



Image gradient

Image gradient

- An image gradient is a directional change in the intensity or color in an image
- For each pixel in the image: G_x , G_y
- ➔ Form a gradient vector (G_x, G_y) :
 - **Important information** to describe the image content
 - Gradient Magnitude = $\sqrt{(G_x)^2 + (G_y)^2} \approx |G_x| + |G_y|$
 - Gradient Direction = $\arctan(G_y/G_x)$



Blue lines represent the gradient direction: from brightest to darkest

Laplacian filter - Second derivative

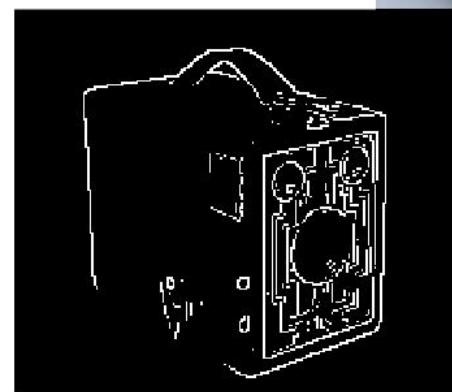
- Discrete approximations for the Laplacian function

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

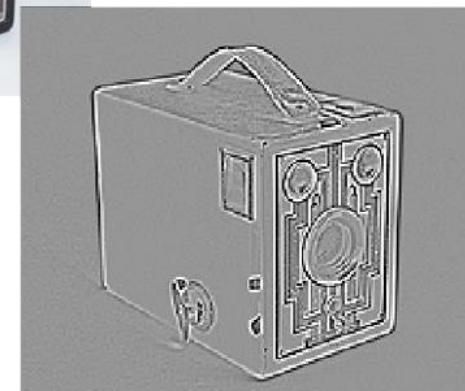
$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- One convolution matrix

Gradient



Laplacian

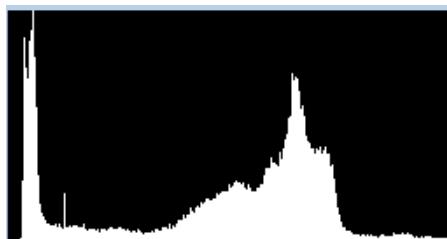


Feature extraction

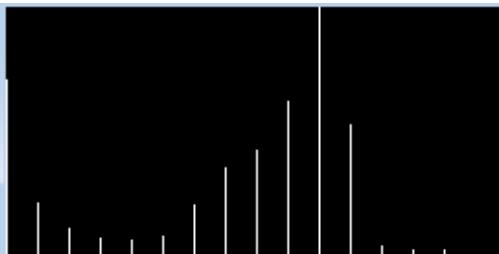
- Two types of features are extracted from the image:
 - local and global features (descriptors)
- Global features
 - Describe the image as a whole to the generalize the entire object
 - Include contour representations, shape descriptors, and texture features
 - Examples: Invariant Moments (Hu, Zerinke), Histogram Oriented Gradients (HOG), PHOG, and Co-HOG,...
- Local feature:
 - the local features describe the image patches (key points in the image) of an object
 - represents the texture/color in an image patch
 - Examples: SIFT, SURF, LBP, BRISK, MSER and FREAK, ...

Feature extraction

■ Global features

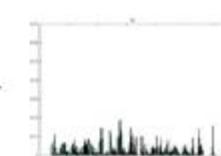
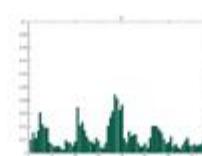
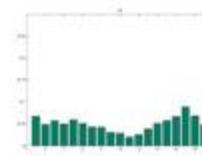
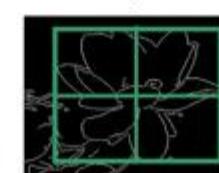


256 bins intensity histogram



16 bins intensity histogram

Input Image (image.jpg)



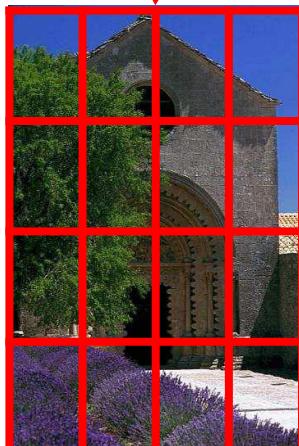
Output PHOG descriptor (image.jpg.txt)

Pyramid Histogram of Oriented Gradients

Feature extraction

- Local features: how to determine image patches / local regions

Dividing into patches with regular grid



Without knowledge about image content

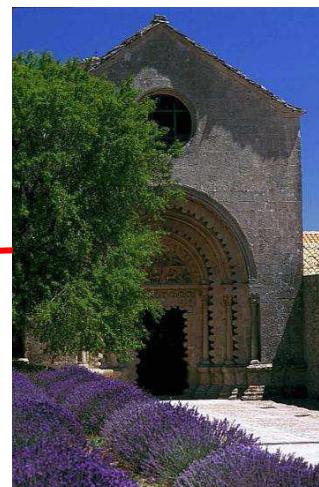


Image segmentation



Keypoint detection



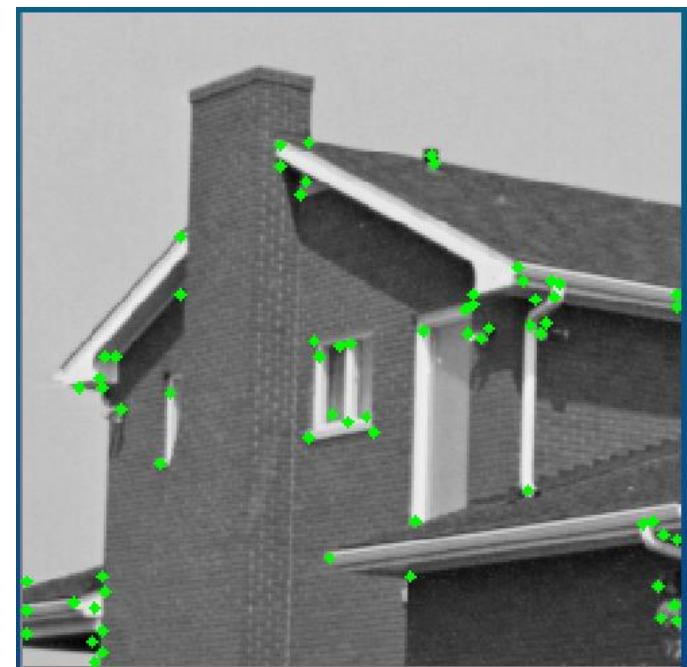
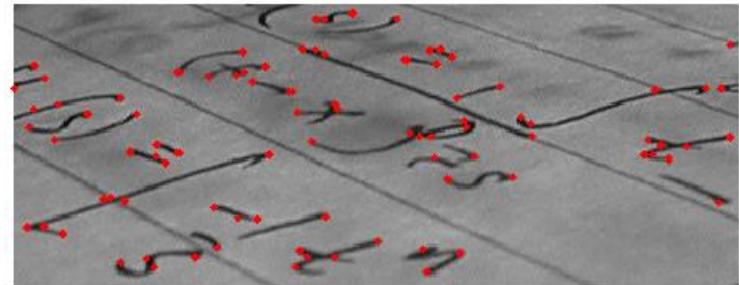
Based on the content of image

Feature extraction

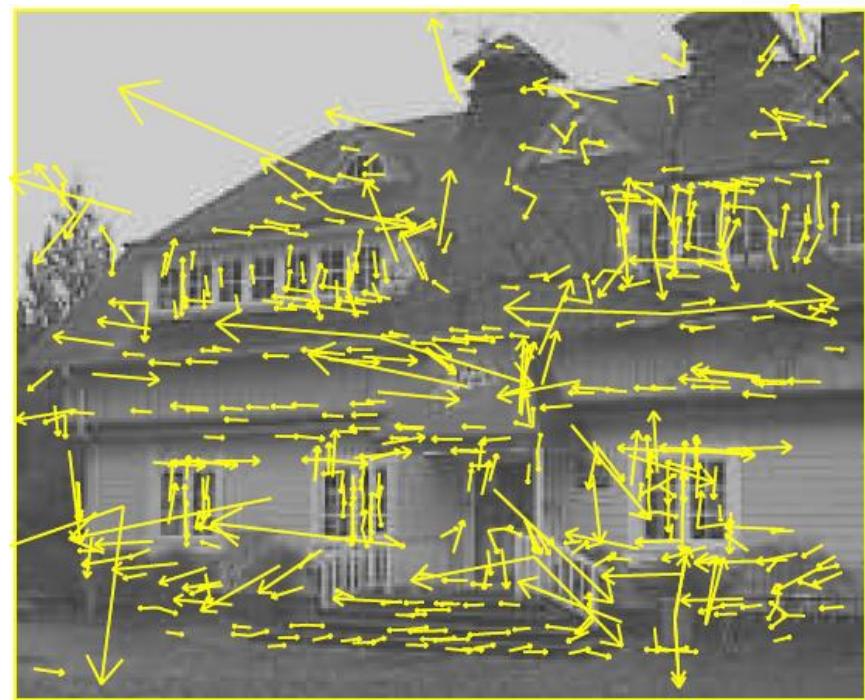
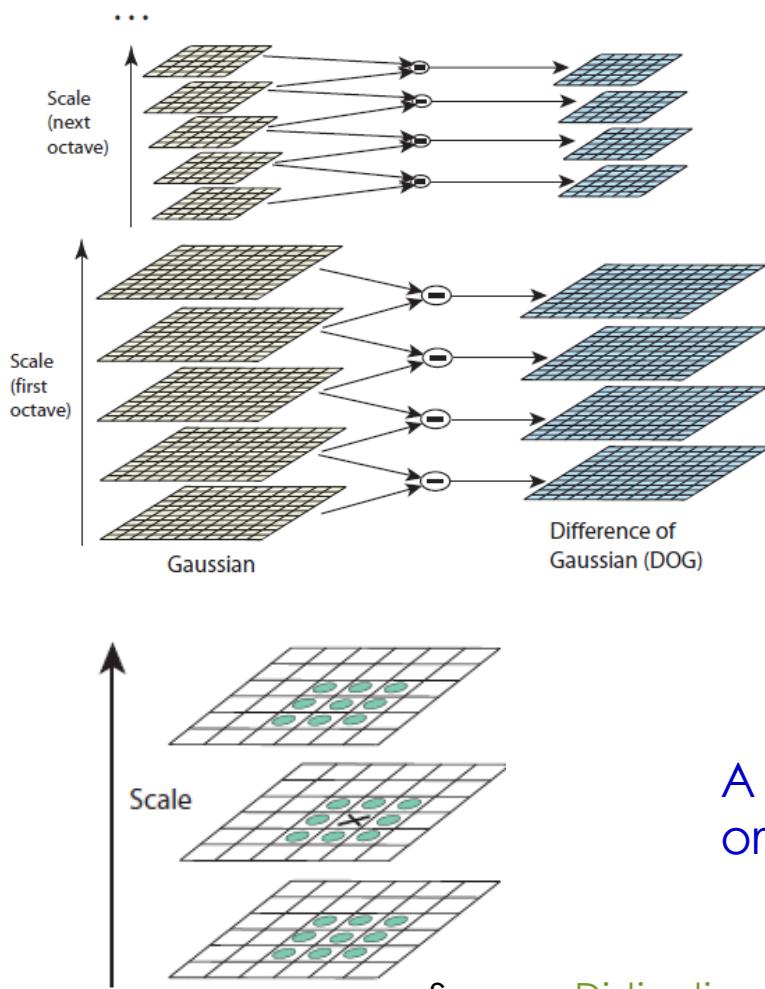
- Image segmentation
 - Thresholding
 - Split and merge
 - Region growing
 - Watershed
 - ...

Feature extraction

- Keypoint detectors:
 - DoG /SIFT detector
 - Harris corner detector
 - Moravec
 - ...
- Local features: computed in local regions associated to each keypoints:
 - SIFT,
 - SURF(Speeded Up Robust Features),
 - PCA-SIFT
 - LBP, BRISK, MSER and FREAK, ...



Feature extraction : DoG/SIFT detector



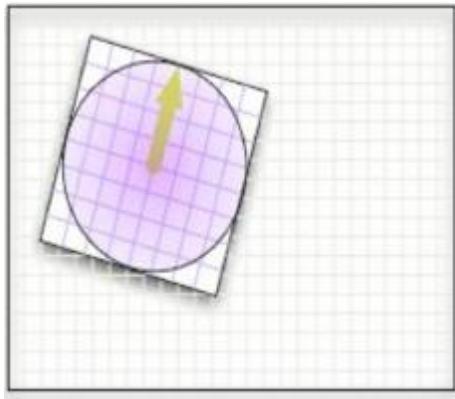
A SIFT keypoint : { $x, y, \text{scale}, \text{dominant orientation}$ }

Source: [Distinctive Image Features from Scale-Invariant Keypoints](#) – IJCV 2004

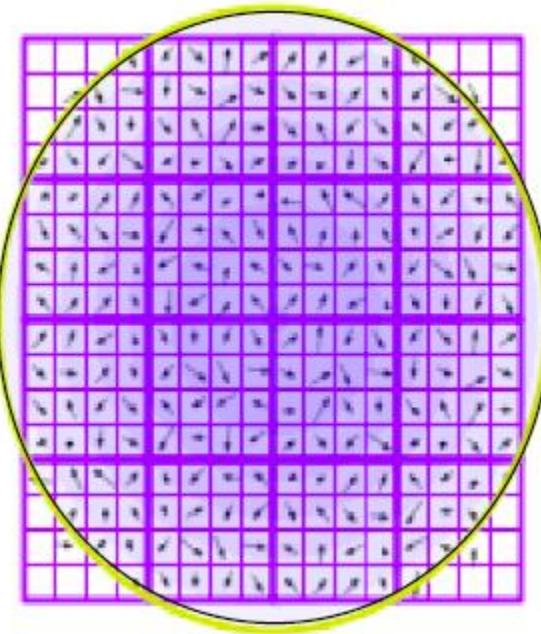
Feature extraction : Good feature?

- Compact
- Invariant to
 - geometric transformation
 - Camera viewpoint
 - Lighting condition
- Best performant local feature: SIFT (David Lowe)

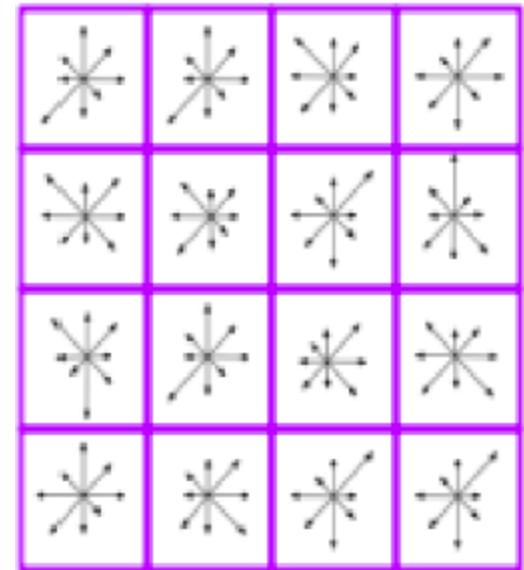
Feature extraction : SIFT feature



Blur the image
using the scale of
the keypoint
(scale invariance)



Compute gradients in
respect to the keypoint
orientation(rotation
invariance)



Compute orientation
histogram in 8
directions over 4x4
sample regions

Other detectors and descriptors

Popular features: SURF, HOG, SIFT

http://campar.in.tum.de/twiki/pub/Chair/TeachingWs13TDCV/feature_descriptors.pdf

Summary some local features:

http://www.cse.iitm.ac.in/~vplab/courses/CV_DIP/PDF/Feature_Detectors_and_Descriptors.pdf

Feature extraction : OpenCV

- SIFT & SURF:
 - Patented algorithms
 - They are free to use fro academic / research purposes
 - You should technically be **getting permission** to use them in **commercial** applications
- From OpenCV 3.0, patented algorithms are
 - removed from standard package,
 - putted into non-free module ([opencv-contrib](#), not installed by default)
- Free alternatives to sif, surf:
 - ORB (Oriented FAST and Rotated Brief)
 - BRIEF, BRISK, FREAK, KAZE and AKAZE

Feature extraction : OpenCV

- SIFT
 - **sift.detect()** function finds the keypoint in the images
 - **sift.compute()** which computes the descriptors from the keypoints

```
sift = cv.xfeatures2d.SIFT_create()
kp = sift.detect(gray,None)
kp,des = sift.compute(gray,kp)
```
 - Find keypoints and descriptors in a single step **sift.detectAndCompute()**

```
sift = cv.xfeatures2d.SIFT_create()
kp, des = sift.detectAndCompute(gray,None)
```
 - https://docs.opencv.org/3.4/da/df5/tutorial_py_sift_intro.html
- SURF: similar

Feature extraction : OpenCV

■ SURF: similar

```
>>> img = cv.imread('fly.png',0)
# Create SURF object. You can specify params here or later.
# Here I set Hessian Threshold to 400
>>> surf = cv.xfeatures2d.SURF_create(400)
# Find keypoints and descriptors directly
>>> kp, des = surf.detectAndCompute(img,None)
>>> len(kp)
699
```



- https://docs.opencv.org/3.4/d7/dd2/tutorial_py_surf_intro.html

References

- CVIP tool to explore the power of computer processing of digital images: Many methods in image processing and computer vision have been implemented
 - <https://cviptools.ece.siu.edu/>
- Library: OpenCV, with C/C++, Python and Java interfaces. OpenCV was designed for computational efficiency and with a strong focus on real-time application:
<https://opencv.org/>
- Books:
 - Rafael C. Gonzalez, Richard Eugene Woods, Digital Image Processing, 2nd edition, Prentice-Hall, 2002: Chap 3 (spatial operators), 6 (Color spaces)
 - Richard Szeliski, Computer Vision: Algorithms and Applications, Springer, 2010.
<http://szeliski.org/Book/>
- Articles:
 - SIFT (DoG detector and SIFT descriptor): <https://www.cs.ubc.ca/~lowe/keypoints/>
 - SURF: Herbert Bay, Andreas Ess, Tinne Tuytelaars, and Luc Van Gool, "[Speeded Up Robust Features](#)", ETH Zurich, Katholieke Universiteit Leuven
 - GLOH: [Krystian Mikolajczyk and Cordelia Schmid "A performance evaluation of local descriptors"](#), IEEE Transactions on Pattern Analysis and Machine Intelligence, 10, 27, pp 1615--1630, 2005.
 - PHOG: <http://www.robots.ox.ac.uk/~vgg/research/caltech/phog.html>
- <https://www.learnopencv.com/> : many examples with code in C++/ Python and clear explanation